

# Program Optimization Through Loop Vectorization

María Garzarán, Saeed Maleki

William Gropp and David Padua

*Department of Computer Science  
University of Illinois at Urbana-Champaign*



# Topics covered

- What are the microprocessor vector extensions or SIMD (Single Instruction Multiple Data Units)
- How to use them
  - Through the compiler via automatic vectorization
    - Manual transformations that enable vectorization
    - Directives to guide the compiler
  - Through intrinsics
- Main focus on vectorizing through the compiler.
  - Code more readable
  - Code portable



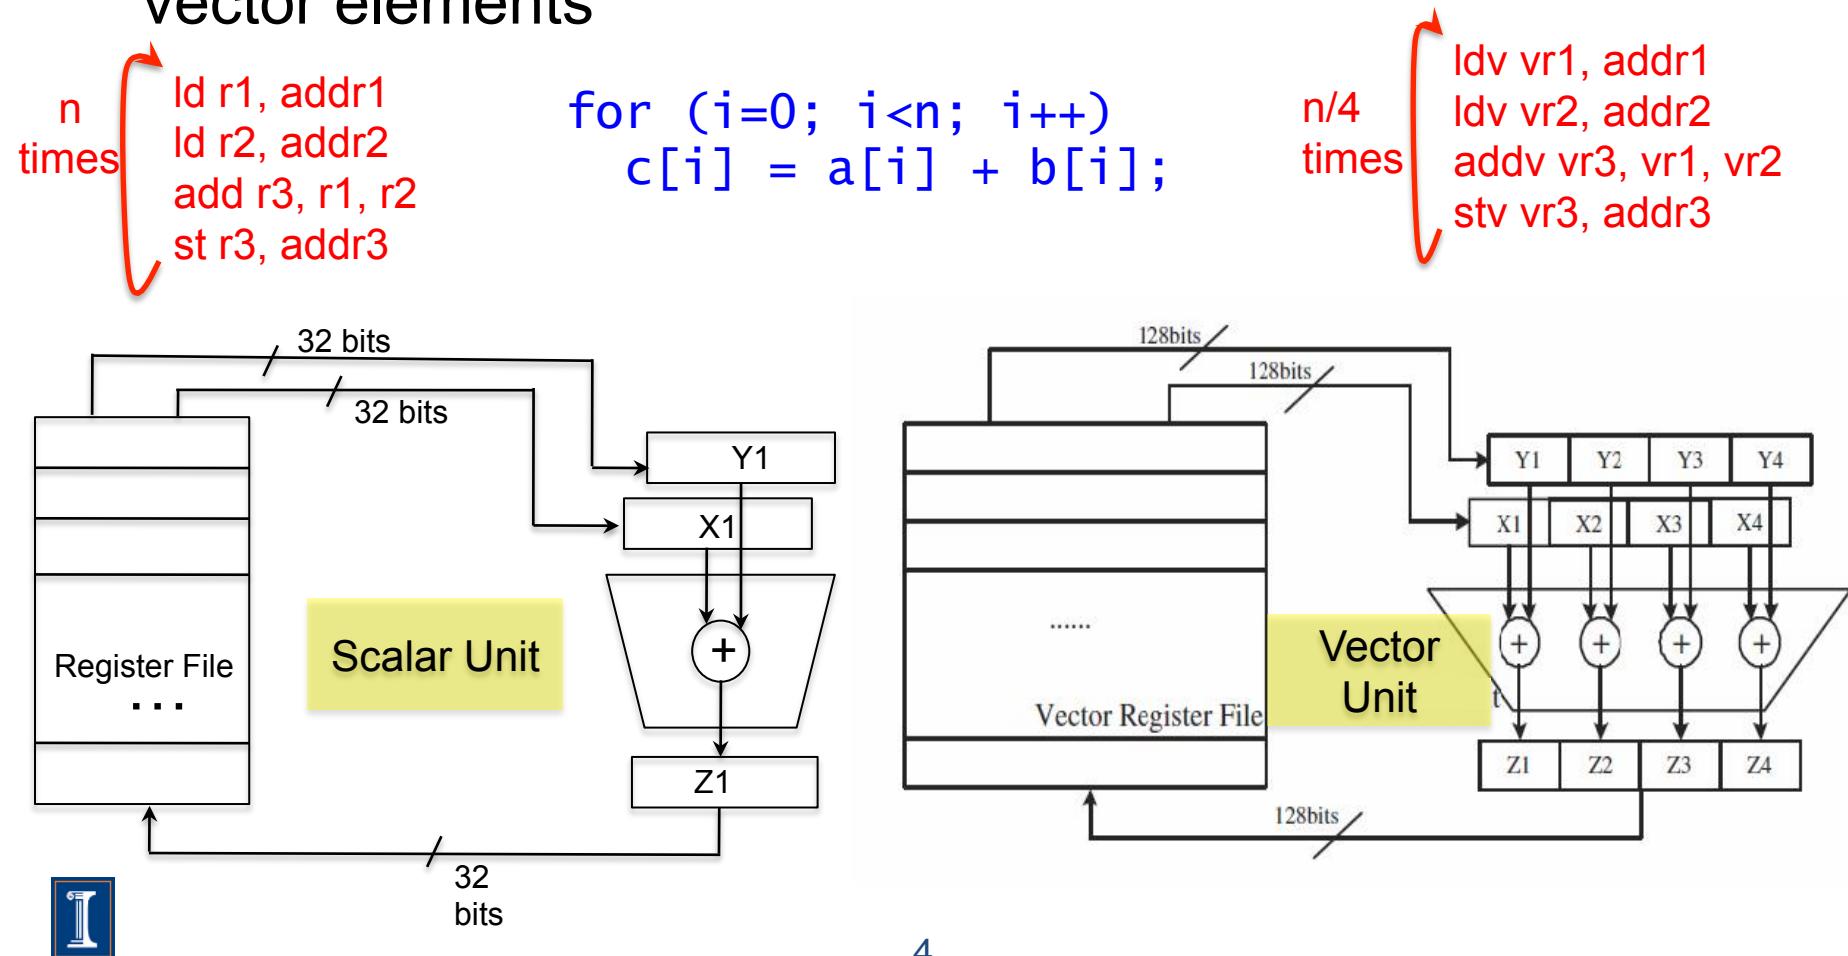
# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
  - Data Alignment
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics



# Simple Example

- Loop vectorization transforms a program so that the same operation is performed at the same time on several vector elements



# SIMD Vectorization

- The use of SIMD units can speed up the program.
- Intel SSE and IBM Altivec have 128-bit vector registers and functional units
  - 4 32-bit single precision floating point numbers
  - 2 64-bit double precision floating point numbers
  - 4 32-bit integer numbers
  - 2 64 bit integer
  - 8 16-bit integer or shorts
  - 16 8-bit bytes or chars
- Assuming a single ALU, these SIMD units can execute 4 single precision floating point number or 2 double precision operations in the time it takes to do only one of these operations by a scalar unit.



# Executing Our Simple Example

S000

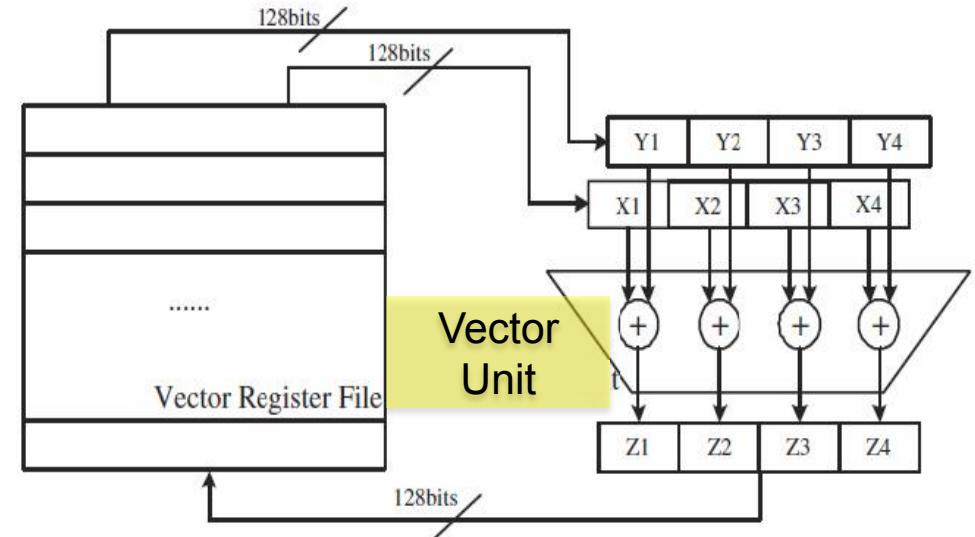
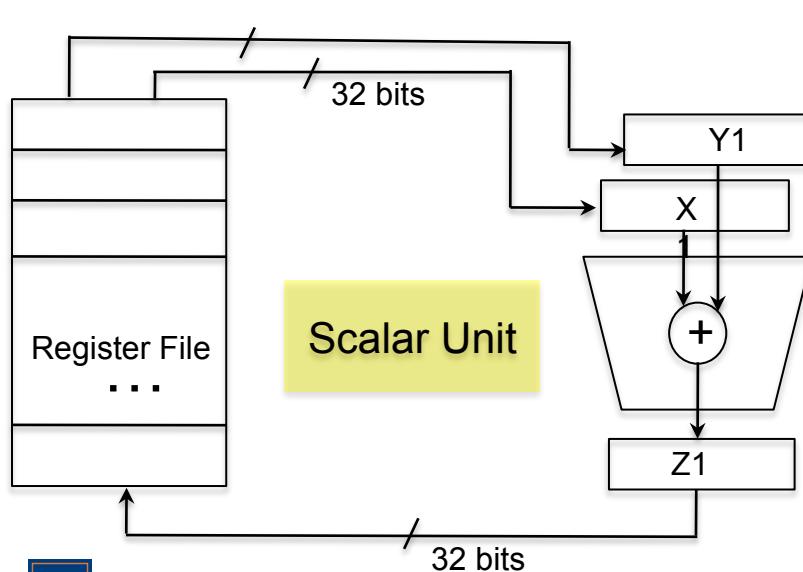
```
for (i=0; i<n; i++)
    c[i] = a[i] + b[i];
```

**Intel Nehalem**

Exec. Time scalar code: 6.1  
Exec. Time vector code: 3.2  
Speedup: 1.8

**IBM Power 7**

Exec. Time scalar code: 2.1  
Exec. Time vector code: 1.0  
Speedup: 2.1



# How do we access the SIMD units?

- Three choices
  1. C code and a vectorizing compiler
  2. Macros or Vector Intrinsics
  3. Assembly Language

```
for (i=0; i<LEN; i++)
    c[i] = a[i] + b[i];
```

```
void example(){
    __m128 rA, rB, rC;
    for (int i = 0; i < LEN; i+=4){
        rA = _mm_load_ps(&a[i]);
        rB = _mm_load_ps(&b[i]);
        rC = _mm_add_ps(rA, rB);
        _mm_store_ps(&c[i], rC);
    }
}
```

```
..B8.5
movaps    a(,%rdx,4), %xmm0
addps    b(,%rdx,4), %xmm0
movaps    %xmm0, c(,%rdx,4)
addq      $4, %rdx
cmpq      $rdi, %rdx
j1       ..B8.5
```



# Why should the compiler vectorize?

1. Easier
2. Portable across vendors and machines
  - Although compiler directives differ across compilers
3. Better performance of the compiler generated code
  - Compiler applies other transformations

Compilers make your codes (almost) machine independent

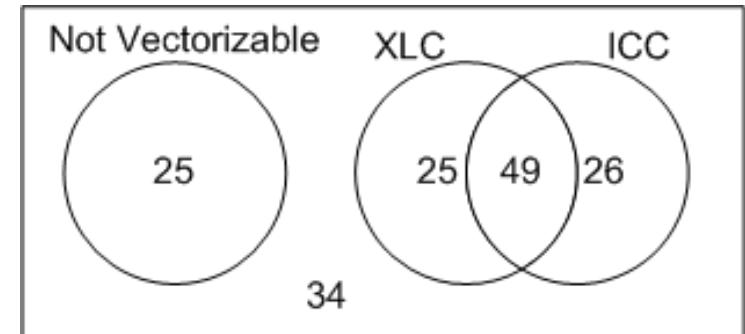
But, compilers fail:

- Programmers need to provide the necessary information
- Programmers need to transform the code



# How well do compilers vectorize?

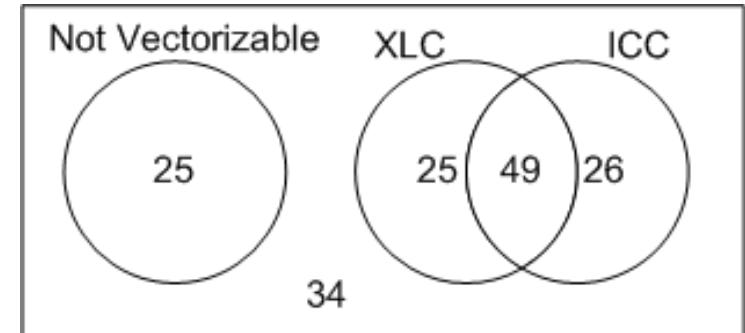
Loops	Compiler	XLC	ICC	GCC
Total		159		
Vectorized		74	75	32
Not vectorized		85	84	127
Average Speed Up		1.73	1.85	1.30



Loops	Compiler	XLC but not ICC	ICC but not XLC
Vectorized		25	26

# How well do compilers vectorize?

Loops	Compiler	XLC	ICC	GCC
Total		159		
Vectorized		74	75	32
Not vectorized		85	84	127
Average Speed Up		1.73	1.85	1.30



Loops	Compiler	XLC but not ICC	ICC but not XLC
Vectorized		25	26

By adding manual vectorization the average speedup was 3.78 (versus 1.73 obtained by the XLC compiler)

# How much programmer intervention?

- Next, three examples to illustrate what the programmer may need to do:
  - Add compiler directives
  - Transform the code
  - Program using vector intrinsics



# Experimental results

- The tutorial shows results for two different platforms with their compilers:
  - Report generated by the compiler
  - Execution Time for each platform

Platform 1: Intel Nehalem  
Intel Core i7 CPU 920@2.67GHz  
Intel ICC compiler, version 11.1  
OS Ubuntu Linux 9.04

Platform 2: IBM Power 7  
IBM Power 7, 3.55 GHz  
IBM xlc compiler, version 11.0  
OS Red Hat Linux Enterprise 5.4

The examples use single precision floating point numbers



# Compiler directives

```
void test(float* A, float* B, float* C, float* D, float* E)
{
    for (int i = 0; i < LEN; i++){
        A[i]=B[i]+C[i]+D[i]+E[i];
    }
}
```



# Compiler directives

S1111

```
void test(float* A, float* B, float*  
C, float* D, float* E)  
{  
    for (int i = 0; i < LEN; i++){  
        A[i]=B[i]+C[i]+D[i]+E[i];  
    }  
}
```

S1111

**Intel Nehalem**  
**Compiler report:** Loop was not  
vectorized.  
**Exec. Time scalar code:** 5.6  
**Exec. Time vector code:** --  
**Speedup:** --



S1111

```
void test(float* __restrict__ A,  
float* __restrict__ B,  
float* __restrict__ C,  
float* __restrict__ D,  
float* __restrict__ E)  
{  
    for (int i = 0; i < LEN; i++){  
        A[i]=B[i]+C[i]+D[i]+E[i];  
    }  
}
```

S1111

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 5.6  
**Exec. Time vector code:** 2.2  
**Speedup:** 2.5

# Compiler directives

S1111

```
void test(float* A, float* B, float*  
C, float* D, float* E)  
{  
    for (int i = 0; i < LEN; i++){  
        A[i]=B[i]+C[i]+D[i]+E[i];  
    }  
}
```

S1111

**Power 7**  
**Compiler report:** Loop was not  
vectorized.  
**Exec. Time scalar code:** 2.3  
**Exec. Time vector code:** --  
**Speedup:** --



S1111

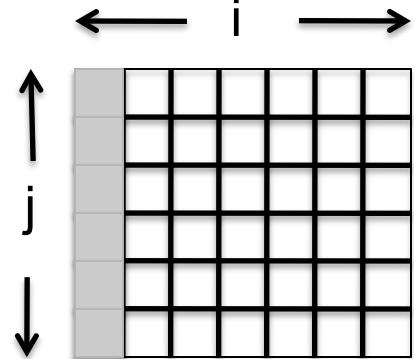
```
void test(float* __restrict__ A,  
float* __restrict__ B,  
float* __restrict__ C,  
float* __restrict__ D,  
float* __restrict__ E)  
{  
    for (int i = 0; i < LEN; i++){  
        A[i]=B[i]+C[i]+D[i]+E[i];  
    }  
}
```

S1111

**Power 7**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 1.6  
**Exec. Time vector code:** 0.6  
**Speedup:** 2.7

# Loop Transformations

```
for (int i=0;i<LEN;i++){  
    sum = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum += A[j][i];  
    }  
    B[i] = sum;  
}
```



```
for (int i=0;i<size;i++){  
    sum[i] = 0;  
    for (int j=0;j<size;j++){  
        sum[i] += A[j][i];  
    }  
    B[i] = sum[i];  
}
```



# Loop Transformations

S136

```
for (int i=0;i<LEN;i++){  
    sum = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum += A[j][i];  
    }  
    B[i] = sum;  
}
```

S136\_1

```
for (int i=0;i<LEN;i++)  
    sum[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum[i] += A[j][i];  
    }  
    B[i]=sum[i];  
}
```

S136\_2

```
for (int i=0;i<LEN;i++)  
    B[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        B[i] += A[j][i];  
    }
```

S136

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Vectorization possible but seems inefficient  
**Exec. Time scalar code:** 3.7  
**Exec. Time vector code:** --  
**Speedup:** --



S136\_1

**Intel Nehalem**  
**report:** Permuted loop was vectorized.  
**scalar code:** 1.6  
**vector code:** 0.6  
**Speedup:** 2.6

S136\_2

**Intel Nehalem**  
**report:** Permuted loop was vectorized.  
**scalar code:** 1.6  
**vector code:** 0.6  
**Speedup:** 2.6

# Loop Transformations

S136

```
for (int i=0;i<LEN;i++){  
    sum = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum += A[j][i];  
    }  
    B[i] = sum;  
}
```

S136\_1

```
for (int i=0;i<LEN;i++)  
    sum[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum[i] += A[j][i];  
    }  
    B[i]=sum[i];  
}
```

S136\_2

```
for (int i=0;i<LEN;i++)  
    B[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        B[i] += A[j][i];  
    }  
}
```

S136

**IBM Power 7**  
**Compiler report:** Loop was not SIMD vectorized  
**Exec. Time scalar code:** 2.0  
**Exec. Time vector code:** --  
**Speedup:** --

S136\_1

**IBM Power 7**  
**report:** Loop interchanging applied.  
Loop was SIMD vectorized  
**scalar code:** 0.4  
**vector code:** 0.2  
**Speedup:** 2.0

S136\_2

**IBM Power 7**  
**report:** Loop interchanging applied.  
Loop was SIMD  
**scalar code:** 0.4  
**vector code:** 0.16  
**Speedup:** 2.7



# Intrinsics (SSE)

```
#define n 1024
__attribute__((aligned(16))) float a[n], b[n], c[n];
int main() {
    for (i = 0; i < n; i++) {
        c[i]=a[i]*b[i];
    }
}
```



```
#include <xmmmintrin.h>
#define n 1024
__attribute__((aligned(16))) float a[n], b[n], c[n];

int main() {
    __m128 rA, rB, rC;
    for (i = 0; i < n; i+=4) {
        rA = _mm_load_ps(&a[i]);
        rB = _mm_load_ps(&b[i]);
        rC= _mm_mul_ps(rA,rB);
        _mm_store_ps(&c[i], rC);
    }
}
```



# Intrinsics (Altivec)

```
#define n 1024
__attribute__ ((aligned(16))) float a[n],b[n],c[n];
...
for (int i=0; i<LEN; i++)
    c[i]=a[i]*b[i];
```



```
vector float rA,rB,rC,r0;          // Declares vector registers
r0 = vec_xor(r0,r0);                // Sets r0 to {0,0,0,0}
for (int i=0; i<LEN; i+=4){        // Loop stride is 4
    rA = vec_ld(0, &a[i]);         // Load values to rA
    rB = vec_ld(0, &b[i]);         // Load values to rB
    rC = vec_madd(rA,rB,r0);      // rA and rB are multiplied
    vec_st(rC, 0, &c[i]);         // rC is stored to the c[i:i+3]
}
```



# Outline

1. Intro
2. Data Dependencies (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependencies
  - Data Alignment
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics



# Data dependences

- The notion of dependence is the foundation of the process of vectorization.
- It is used to build a calculus of program transformations that can be applied manually by the programmer or automatically by a compiler.



# Definition of Dependence

- A statement S is said to be data dependent on statement T if
  - T executes before S in the original sequential/scalar program
  - S and T access the same data item
  - At least one of the accesses is a write.



# Data dependences and vectorization

- Loop dependences guide vectorization
- **Main idea:** A statement inside a loop which is not in a cycle of the dependence graph can be vectorized.

```
for (i=0; i<n; i++){  
S1  a[i] = b[i] + 1;    → a[0:n-1] = b[0:n-1] + 1;  
}
```

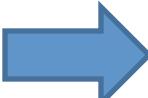
S1



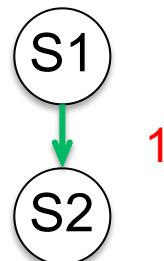
# Data dependences and vectorization

- **Main idea:** A statement inside a loop which is not in a cycle of the dependence graph can be vectorized.

```
for (i=1; i<n; i++){  
S1  a[i] = b[i] + 1;  
S2  c[i] = a[i-1] + 2;  
}
```



```
a[1:n] = b[1:n] + 1;  
c[1:n] = a[0:n-1] + 2;
```



# Data dependences and transformations

- When cycles are present, vectorization can be achieved by:
  - Separating (distributing) the statements not in a cycle
  - Removing dependences
  - Freezing loops
  - Changing the algorithm

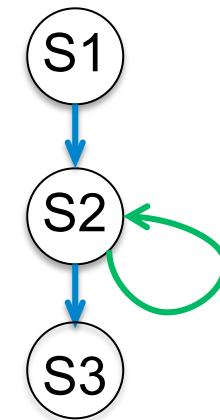


# Distributing

```
for (i=1; i<n; i++){
S1 b[i] = b[i] + c[i];
S2 a[i] = a[i-1]*a[i-2]+b[i];
S3 c[i] = a[i] + 1;
}
```



```
b[1:n-1] = b[1:n-1] + c[1:n-1];
for (i=1; i<n; i++){
    a[i] = a[i-1]*a[i-2]+b[i];
}
c[1:n-1] = a[1:n-1] + 1;
```



# Removing dependences

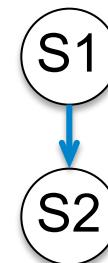
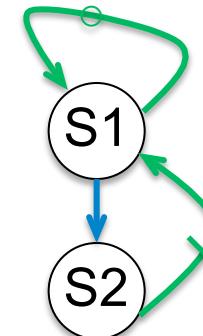
```
for (i=0; i<n; i++){  
S1    a = b[i] + 1;  
S2    c[i] = a + 2;  
}
```



```
for (i=0; i<n; i++){  
S1    a'[i] = b[i] + 1;  
S2    c[i] = a'[i] + 2;  
}  
a=a'[n-1]
```

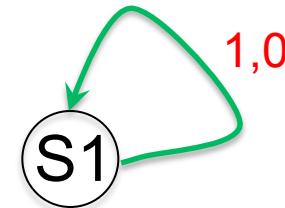


```
S1    a'[0:n-1] = b[0:n-1] + 1;  
S2    c[0:n-1] = a'[0:n-1] + 2;  
a=a'[n-1]
```



# Freezing Loops

```
for (i=1; i<n; i++) {  
    for (j=1; j<n; j++) {  
        a[i][j]=a[i][j]+a[i-1][j];  
    }  
}
```



Ignoring (freezing) the outer loop:

```
for (j=1; j<n; j++) {  
    a[i][j]=a[i][j]+a[i-1][j];  
}
```



```
for (i=1; i<n; i++) {  
    a[i][1:n-1]=a[i][1:n-1]+a[i-1][1:n-1];  
}
```



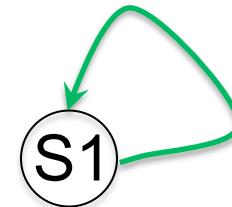
# Changing the algorithm

- When there is a recurrence, it is necessary to change the algorithm in order to vectorize.
- Compiler use pattern matching to identify the recurrence and then replace it with a parallel version.
- Examples of recurrences include:
  - Reductions (`S+=A[i]`)
  - Linear recurrences (`A[i]=B[i]*A[i-1]+C[i]`)
  - Boolean recurrences (`if (A[i]>max) max = A[i]`)



# Changing the algorithm (cont.)

```
S1  a[0]=b[0];  
    for (i=1; i<n; i++)  
S2      a[i]=a[i-1]+b[i];
```



```
a[0:n-1]=b[0:n-1];  
for (i=0;i<k;i++) /* n = 2k */  
    a[2**i:n-1]=a[2**i:n-1]+b[0:n-2**i];
```



# Stripmining

- Stripmining is a simple transformation.

```
for (i=1; i<n; i++){           /* n is a multiple of q */  
}    ...      for (k=1; k<n; k+=q){  
        for (i=k; i<k+q-1; i++){  
            ...  
        }  
    }  
}
```

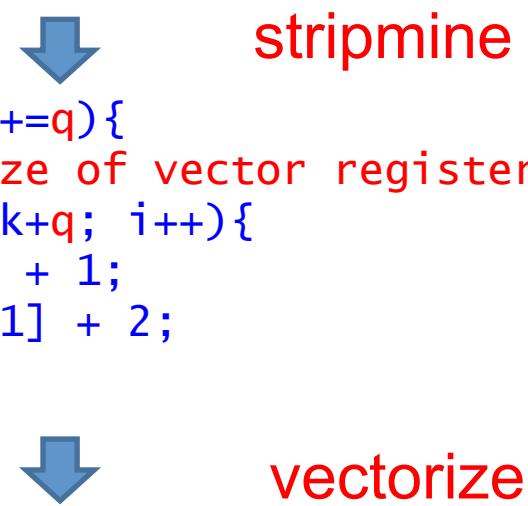
- It is typically used to improve locality.



# Stripmining (cont.)

- Stripmining is often used when vectorizing

```
for (i=1; i<n; i++){  
    a[i] = b[i] + 1;  
    c[i] = a[i] + 2;  
}  
  
for (k=1; k<n; k+=q){  
    /* q could be size of vector register */  
    for (i=k; i < k+q; i++){  
        a[i] = b[i] + 1;  
        c[i] = a[i-1] + 2;  
    }  
}  
  
for (i=1; i<n; i+=q){  
    a[i:i+q-1] = b[i:i+q-1] + 1;  
    c[i:i+q-1] = a[i:i+q] + 2;  
}
```



# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
  - Data Alignment
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics



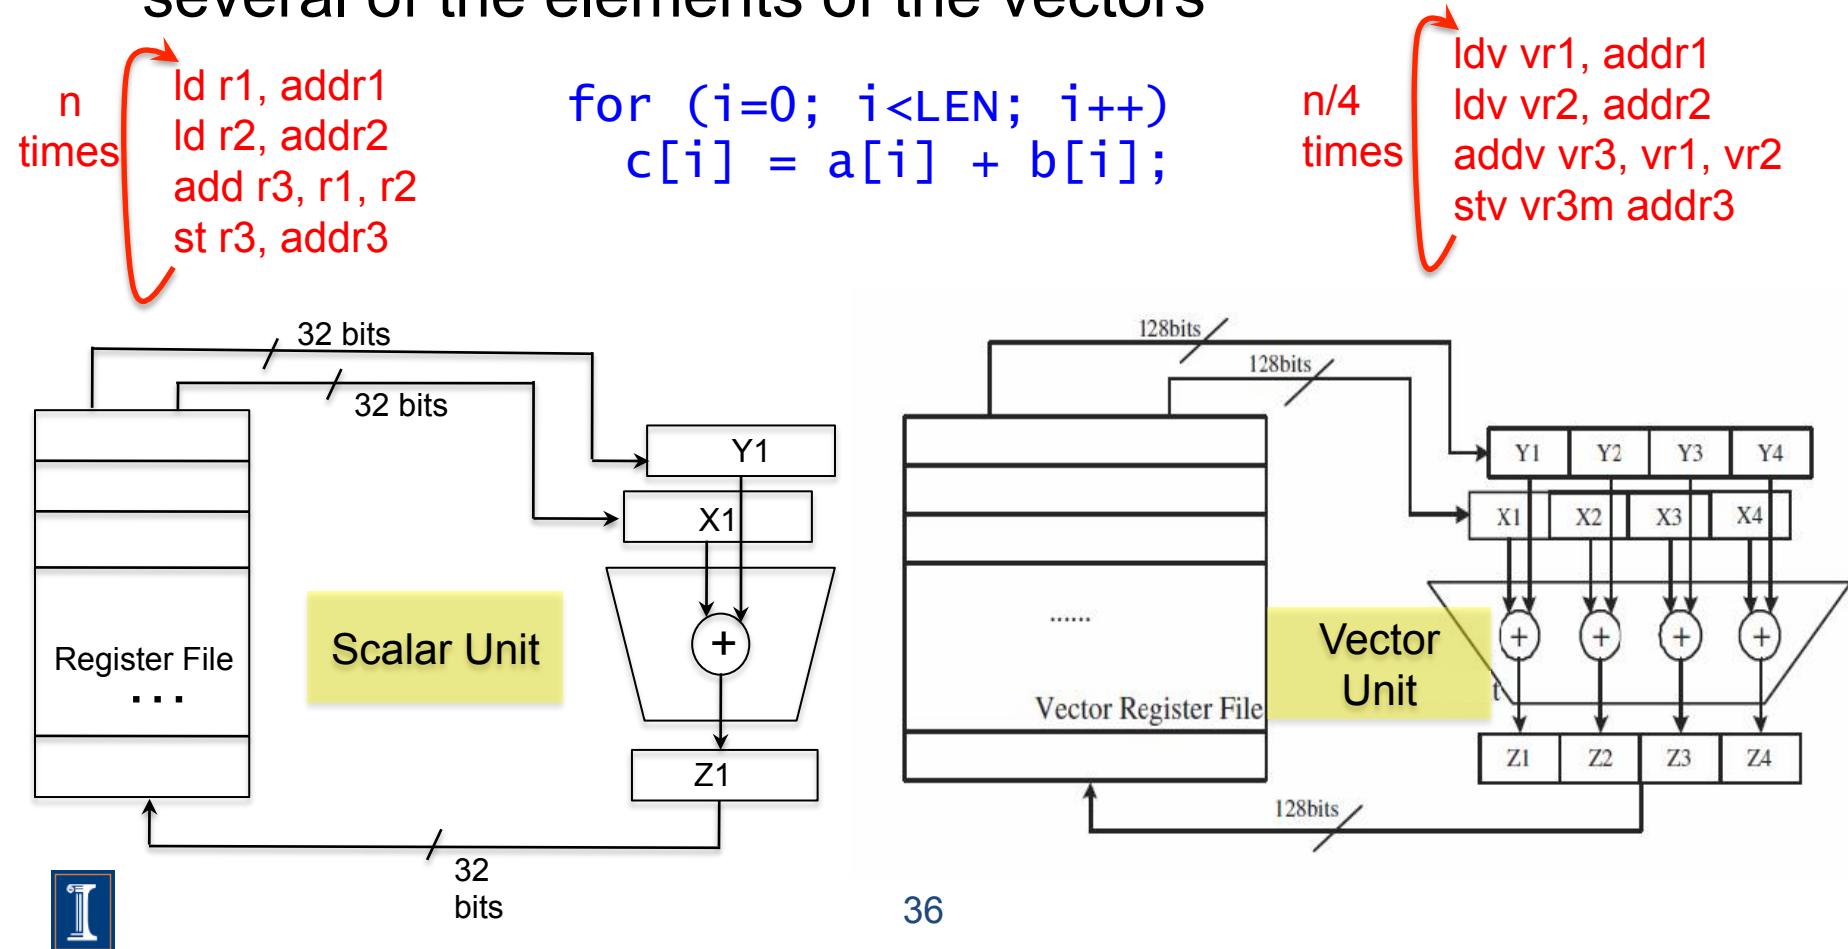
# Loop Vectorization

- Loop Vectorization is not always a legal and profitable transformation.
- Compiler needs:
  - Compute the dependences
    - The compiler figures out dependences by
      - Solving a system of (integer) equations (with constraints)
      - Demonstrating that there is no solution to the system of equations
    - Remove cycles in the dependence graph
    - Determine data alignment
    - Vectorization is profitable



# Simple Example

- Loop vectorization transforms a program so that the same operation is performed at the same time on several of the elements of the vectors



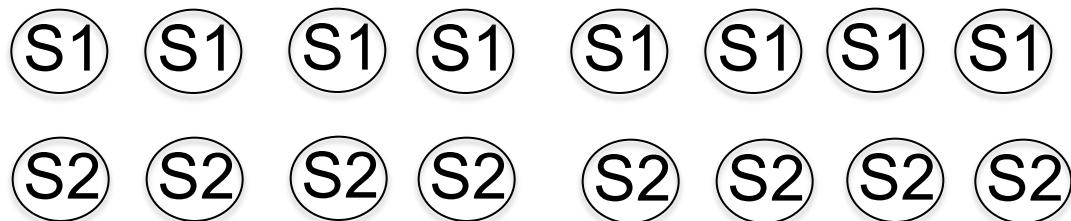
# Loop Vectorization

- When vectorizing a loop with several statements the compiler need to strip-mine the loop and then apply loop distribution

```
for (i=0; i<LEN; i++){  
S1 a[i]=b[i]+(float)1.0;  
S2 c[i]=b[i]+(float)2.0;  
}
```

```
for (i=0; i<LEN; i+=strip_size){  
    for (j=i; j<i+strip_size; j++)  
        a[j]=b[j]+(float)1.0;  
    for (j=i; j<i+strip_size; j++)  
        c[j]=b[j]+(float)2.0;  
}
```

i=0   i=1   i=2   i=3   i=4   i=5   i=6   i=7



# Loop Vectorization

- When vectorizing a loop with several statements the compiler need to strip-mine the loop and then apply loop distribution

```
for (i=0; i<LEN; i++){  
S1 a[i]=b[i]+(float)1.0;  
S2 c[i]=b[i]+(float)2.0;  
}
```

```
for (i=0; i<LEN; i+=strip_size){  
    for (j=i; j<i+strip_size; j++)  
        a[j]=b[j]+(float)1.0;  
    for (j=i; j<i+strip_size; j++)  
        c[j]=b[j]+(float)2.0;  
}
```

i=0    i=1    i=2    i=3    i=4    i=5    i=6    i=7



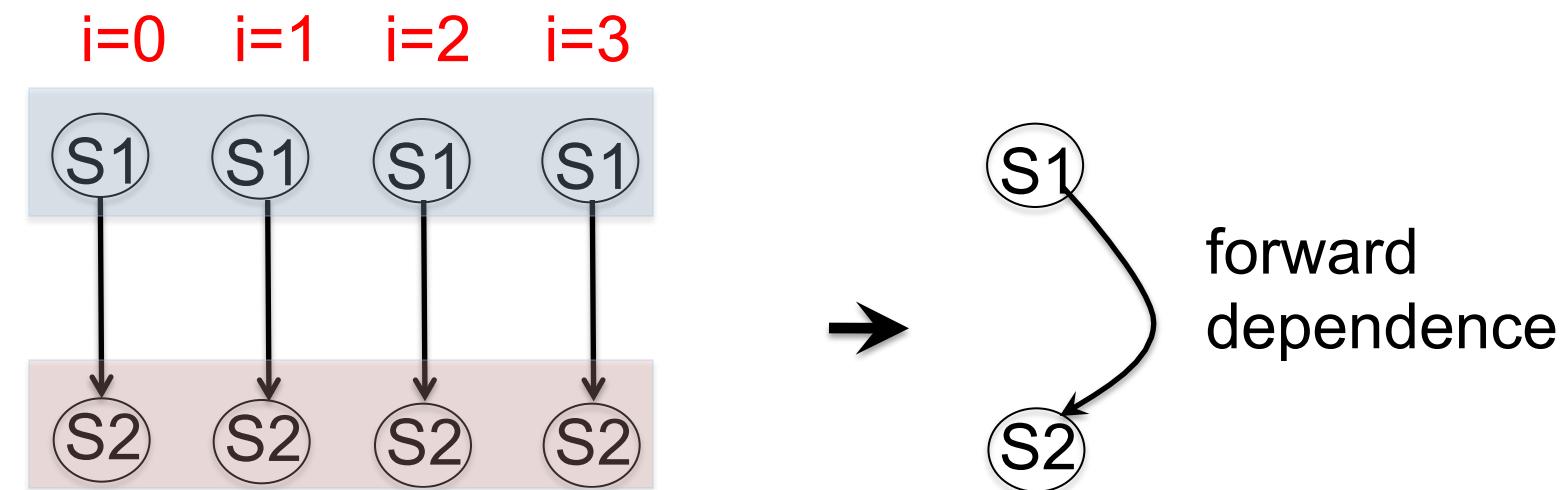
# Dependence Graphs and Compiler Vectorization

- No dependences: previous two slides
- Acyclic graphs:
  - All dependences are forward:
    - Vectorized by the compiler
  - Some backward dependences:
    - Sometimes vectorized by the compiler
- Cycles in the dependence graph
  - Self-antidependence:
    - Vectorized by the compiler
  - Recurrence:
    - Usually not vectorized by the compiler
  - Other examples



# Acyclic Dependence Graphs: Forward Dependencies

```
for (i=0; i<LEN; i++) {  
S1 a[i]= b[i] + c[i]  
S2 d[i] = a[i] + (float) 1.0;  
}
```



# Acyclic Dependence Graphs: Forward Dependencies

S113

```
for (i=0; i<LEN; i++) {  
    a[i] = b[i] + c[i]  
    d[i] = a[i] + (float) 1.0;  
}
```

## Intel Nehalem

**Compiler report:** Loop was  
vectorized

**Exec. Time scalar code:** 10.2

**Exec. Time vector code:** 6.3

**Speedup:** 1.6

## IBM Power 7

**Compiler report:** Loop was SIMD  
vectorized

**Exec. Time scalar code:** 3.1

**Exec. Time vector code:** 1.5

**Speedup:** 2.0



# Acyclic Dependence Graphs

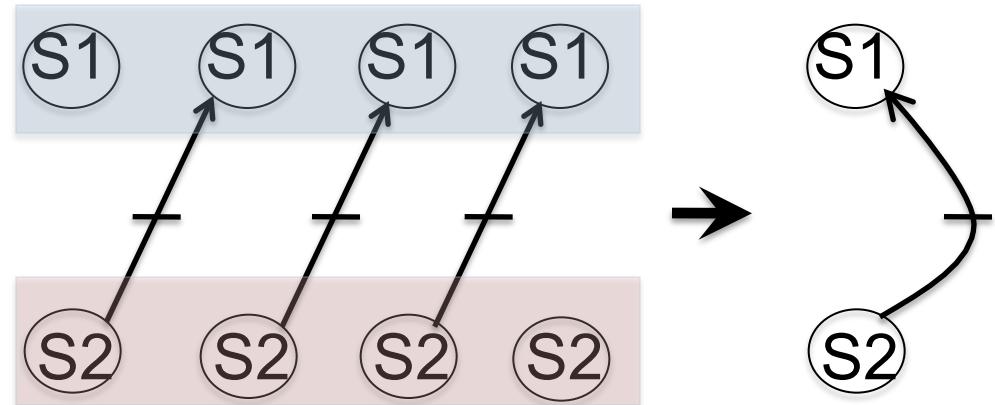
## Backward Dependencies (I)

```
for (i=0; i<LEN; i++) {  
S1  a[i] = b[i] + c[i]  
S2  d[i] = a[i+1] + (float) 1.0;  
}
```

backward  
dependence

i=0  
 $\left\{ \begin{array}{l} S1: a[0] = b[0] + c[0] \\ S2: d[0] = a[1] + 1 \end{array} \right.$

i=1  
 $\left\{ \begin{array}{l} S1: a[1] = b[0] + c[0] \\ S2: d[1] = a[2] + 1 \end{array} \right.$



This loop cannot be vectorized as it is



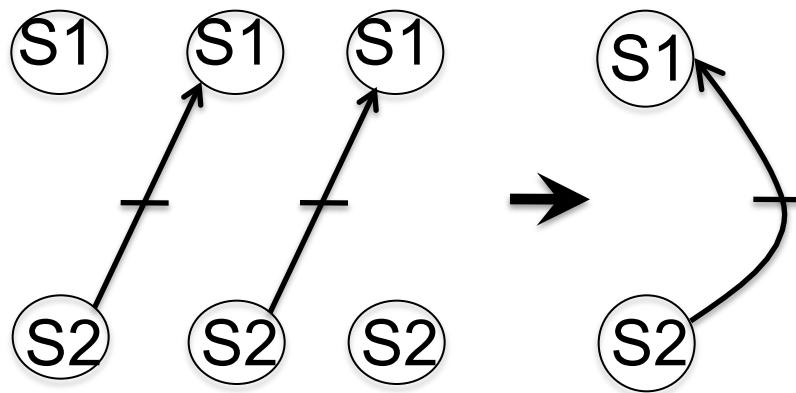
# Acyclic Dependence Graphs

## Backward Dependencies (I)

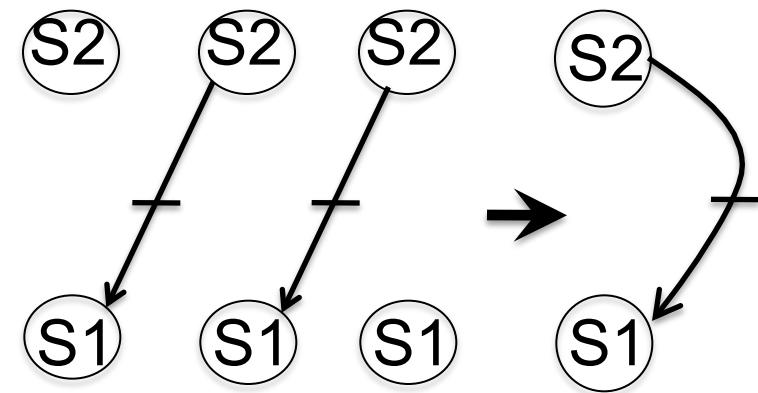
Reorder of statements

```
for (i=0; i<LEN; i++) {  
S1 a[i]= b[i] + c[i]  
S2 d[i] = a[i+1] + (float) 1.0;  
}  
}
```

```
for (i=0; i<LEN; i++) {  
S2 d[i] = a[i+1]+(float)1.0;  
S1 a[i]= b[i] + c[i];  
}  
}
```



backward  
dependedence



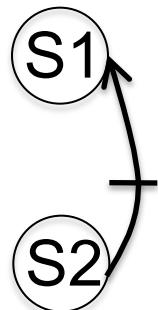
forward  
dependedence



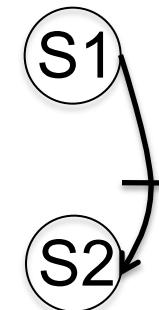
# Acyclic Dependence Graphs

## Backward Dependencies (I)

```
S114  
for (i=0; i<LEN; i++) {  
    a[i] = b[i] + c[i];  
    d[i] = a[i+1]+(float)1.0;  
}
```



```
S114_1  
for (i=0; i<LEN; i++) {  
    d[i] = a[i+1]+(float)1.0;  
    a[i] = b[i] + c[i];  
}
```



**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Existence of vector dependence  
**Exec. Time scalar code:** 12.6  
**Exec. Time vector code:** --  
**Speedup:** --

**Intel Nehalem**  
**Compiler report:** Loop was vectorized  
**Exec. Time scalar code:** 10.7  
**Exec. Time vector code:** 9.4  
**Speedup:** 1.03  
**Speedup vs non-reordered code:** 1.35



# Acyclic Dependence Graphs

## Backward Dependencies (I)

S114

```
for (i=0; i<LEN; i++) {  
    a[i] = b[i] + c[i];  
    d[i] = a[i+1]+(float)1.0;  
}
```

S114\_1

```
for (i=0; i<LEN; i++) {  
    d[i] = a[i+1]+(float)1.0;  
    a[i] = b[i] + c[i];  
}
```

The IBM XLC compiler generated the same code in both cases

S114

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 1.2  
**Exec. Time vector code:** 0.6  
**Speedup:** 2.0

S114\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 1.2  
**Exec. Time vector code:** 0.6  
**Speedup:** 2.0



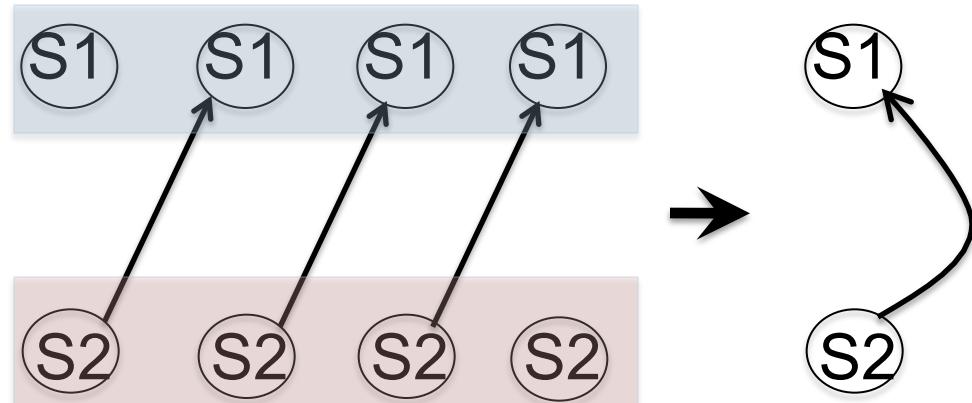
# Acyclic Dependence Graphs

## Backward Dependencies (II)

```
for (int i = 1; i < LEN; i++) {  
S1 a[i] = d[i-1] + (float)sqrt(c[i]);  
S2 d[i] = b[i] + (float)sqrt(e[i]);  
}
```

backward  
dependence

i=1 {  
S1: a[1] = d[0] + sqrt(c[1])  
S2: d[1] = b[1] + sqrt(e[1])  
}  
i=2 {  
S1: a[2] = d[1] + sqrt(c[2])  
S2: d[2] = b[2] + sqrt(e[2])  
}

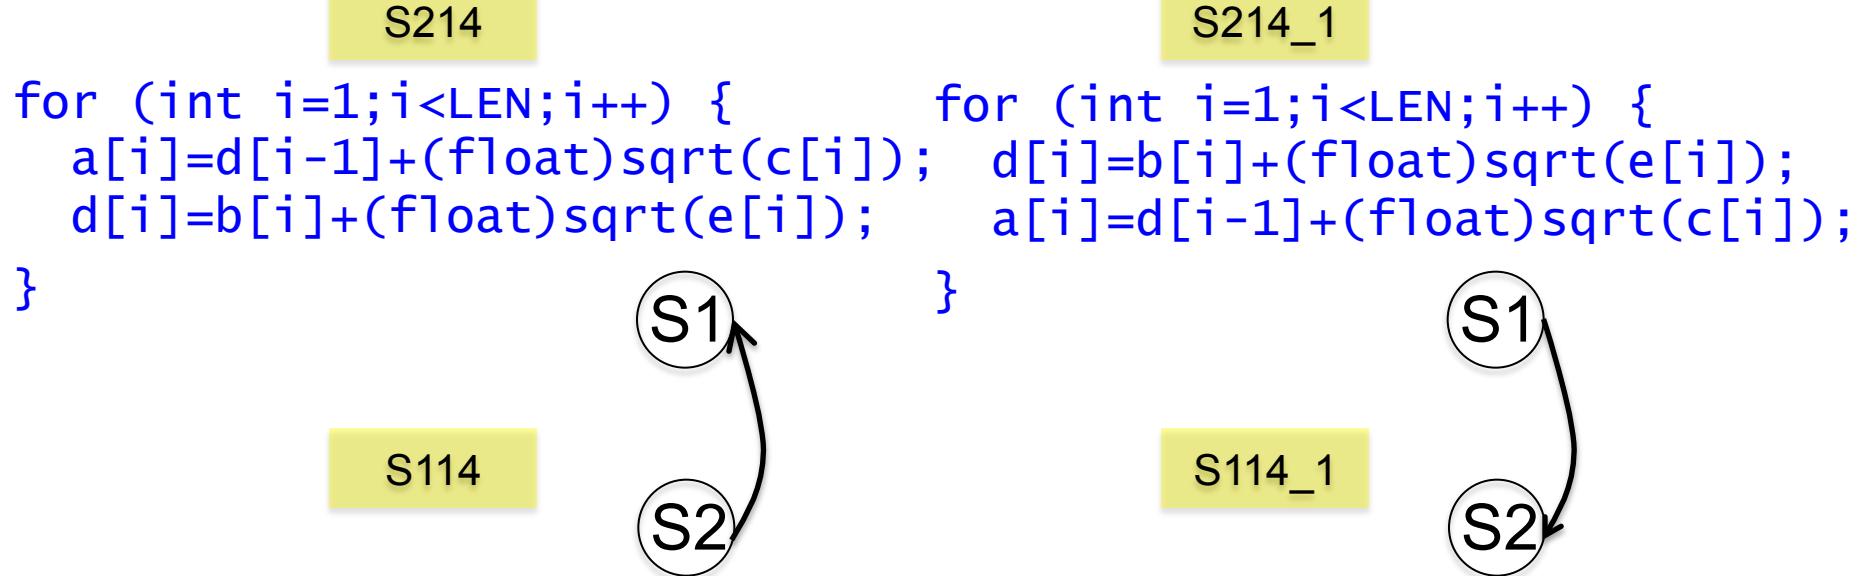


This loop cannot be vectorized as it is



# Acyclic Dependenden Graphs

## Backward Dependencies (II)



### Intel Nehalem

**Compiler report:** Loop was not vectorized. Existence of vector dependence

**Exec. Time scalar code:** 7.6  
**Exec. Time vector code:** --  
**Speedup:** --

### Intel Nehalem

**Compiler report:** Loop was vectorized  
**Exec. Time scalar code:** 7.6  
**Exec. Time vector code:** 3.8  
**Speedup:** 2.0



# Acyclic Dependences Graphs

## Backward Dependencies (II)

S114

```
for (i=0; i<LEN; i++) {  
    a[i] = b[i] + c[i];  
    d[i] = a[i+1]+(float)1.0;  
}
```

S114\_1

```
for (i=0; i<LEN; i++) {  
    d[i] = a[i+1]+(float)1.0;  
    a[i] = b[i] + c[i];  
}
```

The IBM XLC compiler generated the same code in both cases

S114

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 3.3  
**Exec. Time vector code:** 1.8  
**Speedup:** 1.8

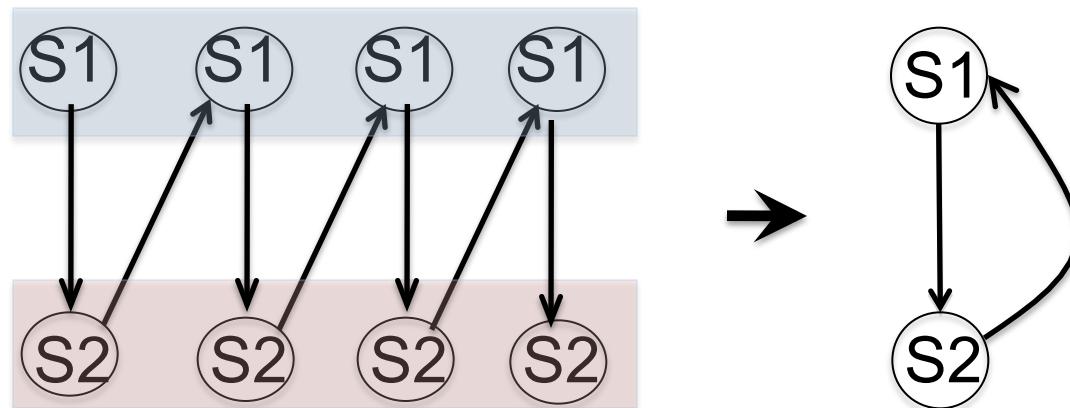
S114\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 3.3  
**Exec. Time vector code:** 1.8  
**Speedup:** 1.8



# Cycles in the DG (I)

```
for (int i=0;i<LEN-1;i++){
S1  b[i] = a[i] + (float) 1.0;
S2  a[i+1] = b[i] + (float) 2.0;
}
```



This loop cannot be vectorized (as it is)  
Statements cannot be simply reordered

# Cycles in the DG (I)

S115

```
for (int i=0;i<LEN-1;i++){
    b[i] = a[i] + (float) 1.0;
    a[i+1] = b[i] + (float) 2.0;
}
```

S115

**Intel Nehalem**

**Compiler report:** Loop was not vectorized.

Existence of vector dependence

**Exec. Time scalar code:** 12.1

**Exec. Time vector code:** --

**Speedup:** --



# Cycles in the DG (I)

S115

```
for (int i=0;i<LEN-1;i++){
    b[i] = a[i] + (float) 1.0;
    a[i+1] = b[i] + (float) 2.0;
}
```

S115

**IBM Power 7**

**Compiler report:** Loop was SIMD vectorized

**Exec. Time scalar code:** 3.1

**Exec. Time vector code:** 2.2

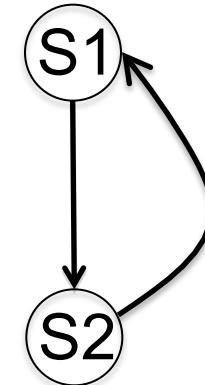
**Speedup:** 1.4



# Cycles in the DG (I)

S115

```
for (int i=0;i<LEN-1;i++){
    b[i] = a[i] + (float) 1.0;
    a[i+1] = b[i] + (float) 2.0;
}
```

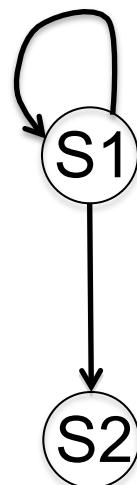


The IBM XLC compiler applies forward substitution and reordering to vectorize the code

compiler generated code

This loop is →  
not vectorized    for (int i=0;i<LEN-1;i++)
 a[i+1]=a[i]+(float)1.0+(float)2.0;

This loop is →  
vectorized      for (int i=0;i<LEN-1;i++)
 b[i] = a[i] + (float) 1.0;



# Cycles in the DG (I)

S115

```
for (int i=0;i<LEN-1;i++){  
    b[i] =a[i]+(float)1.0;  
    a[i+1]=b[i]+(float)2.0;  
}
```

S215

```
for (int i=0;i<LEN-1;i++){  
    b[i]=a[i]+d[i]*d[i]+c[i]*c[i]+c[i]*d[i];  
    a[i+1]=b[i]+(float)2.0;  
}
```

Will the IBM XLC compiler  
vectorize this code as before?



# Cycles in the DG (I)

S115

```
for (int i=0;i<LEN-1;i++){
    b[i] =a[i]+(float)1.0;
    a[i+1]=b[i]+(float)2.0;
}
```

S215

```
for (int i=0;i<LEN-1;i++){
    b[i]=a[i]+d[i]*d[i]+c[i]*c[i]+c[i]*d[i];
    a[i+1]=b[i]+(float)2.0;
}
```

Will the IBM XLC compiler  
vectorize this code as before?

To vectorize, the compiler needs to do this

```
for (int i=0;i<LEN-1;i++)
    a[i+1]=a[i]+d[i]*d[i]+c[i]*c[i]+c[i]*d[i]+(float)2.0;

for (int i=0;i<LEN-1;i++)
    b[i]=a[i]+d[i]*d[i]+c[i]*c[i]+c[i]*d[i]+(float) 1.0;
```



# Cycles in the DG (I)

S115

```
for (int i=0;i<LEN-1;i++){
    b[i] =a[i]+(float)1.0;
    a[i+1]=b[i]+(float)2.0;
}
```

```
for (int i=0;i<LEN-1;i++)
    a[i+1]=a[i]+d[i]*d[i]+c[i]*c[i]+c[i]*d[i]+(float)2.0;
```

```
for (int i=0;i<LEN-1;i++)
    b[i]=a[i]+d[i]*d[i]+c[i]*c[i]+c[i]*d[i]+(float) 1.0;
```

S215

```
for (int i=0;i<LEN-1;i++){
    b[i]=a[i]+d[i]*d[i]+c[i]*c[i]+c[i]*d[i];
    a[i+1]=b[i]+(float)2.0;
}
```

Will the IBM XLC compiler  
vectorize this code as before?

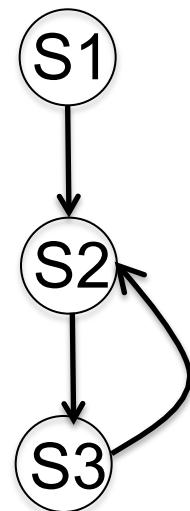
No, the compiler does not  
vectorize S215 because  
it is not cost-effective



# Cycles in the DG (II)

A loop can be partially vectorized

```
for (int i=1;i<LEN;i++){  
S1  a[i] = b[i] + c[i];  
S2  d[i] = a[i] + e[i-1];  
S3  e[i] = d[i] + c[i];  
}
```



S1 can be vectorized  
S2 and S3 cannot be vectorized (as they are)



# Cycles in the DG (II)

S116

```
for (int i=1;i<LEN;i++){  
    a[i] = b[i] + c[i];  
    d[i] = a[i] + e[i-1];  
    e[i] = d[i] + c[i];  
}
```

S116

```
for (int i=1;i<LEN;i++){  
    a[i] = b[i] + c[i];  
    d[i] = a[i] + e[i-1];  
    e[i] = d[i] + c[i];  
}
```

S116

**Intel Nehalem**  
**Compiler report:** Loop was partially vectorized  
**Exec. Time scalar code:** 14.7  
**Exec. Time vector code:** 18.1  
**Speedup:** 0.8

S116

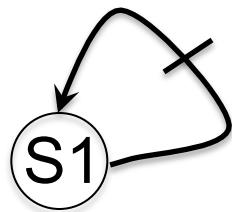
**IBM Power 7**  
**Compiler report:** Loop was not SIMD vectorized because a data dependence prevents SIMD vectorization  
**Exec. Time scalar code:** 13.5  
**Exec. Time vector code:** --  
**Speedup:** --



# Cycles in the DG (III)

```
for (int i=0;i<LEN-1;i++){  
S1  a[i]=a[i+1]+b[i];  
}
```

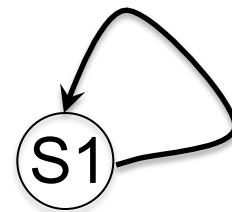
$$\begin{aligned}a[0] &= a[1]+b[0] \\a[1] &= a[2]+b[1] \\a[2] &= a[3]+b[2] \\a[3] &= a[4]+b[3]\end{aligned}$$



Self-antidependence  
can be vectorized

```
for (int i=1;i<LEN;i++){  
S1  a[i]=a[i-1]+b[i];  
}
```

$$\begin{aligned}a[1] &= a[0]+b[1] \\a[2] &= a[1]+b[2] \\a[3] &= a[2]+b[3] \\a[4] &= a[3]+b[4]\end{aligned}$$



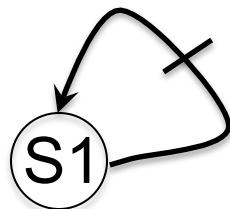
Self true-dependence  
**can not** vectorized  
(as it is)



# Cycles in the DG (III)

S117

```
for (int i=0;i<LEN-1;i++){  
S1 a[i]=a[i+1]+b[i];  
}
```

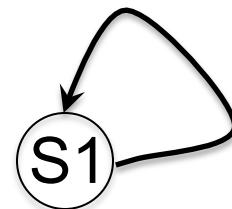


S117

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized  
**Exec. Time scalar code:** 6.0  
**Exec. Time vector code:** 2.7  
**Speedup:** 2.2

S118

```
for (int i=1;i<LEN;i++){  
S1 a[i]=a[i-1]+b[i];  
}
```



S118

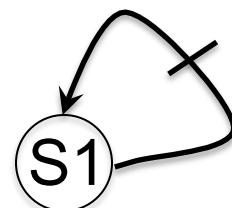
**Intel Nehalem**  
**Compiler report:** Loop was not  
vectorized. Existence of vector  
dependence  
**Exec. Time scalar code:** 7.2  
**Exec. Time vector code:** --  
**Speedup:** --



# Cycles in the DG (III)

S117

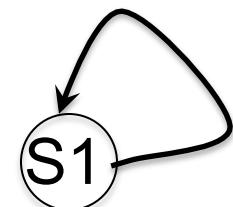
```
for (int i=0;i<LEN-1;i++){  
S1 a[i]=a[i+1]+b[i];  
}
```



S117

S118

```
for (int i=1;i<LEN;i++){  
S1 a[i]=a[i-1]+b[i];  
}
```



S118

## IBM Power 7

**Compiler report:** Loop was SIMD vectorized

**Exec. Time scalar code:** 2.0

**Exec. Time vector code:** 1.0

**Speedup:** 2.0

## IBM Power 7

**Compiler report:** Loop was not SIMD vectorized because a data dependence prevents SIMD vectorization

**Exec. Time scalar code:** 7.2

**Exec. Time vector code:** --

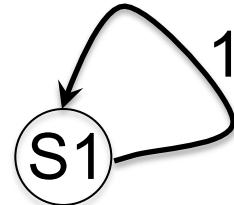
**Speedup:** --



# Cycles in the DG (IV)

```
for (int i=1;i<LEN;i++){  
S1  a[i]=a[i-1]+b[i];  
}
```

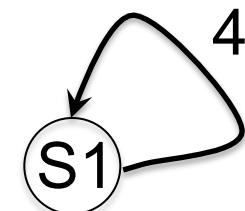
$$\begin{aligned} a[1] &= a[0]+b[1] \\ a[2] &= a[1]+b[2] \\ a[3] &= a[2]+b[3] \end{aligned}$$



Self true-dependence  
is **not** vectorized

```
for (int i=4;i<LEN;i++){  
    a[i]=a[i-4]+b[i];  
}
```

$$\begin{aligned} i=4 \quad a[4] &= a[0]+b[4] \\ i=5 \quad a[5] &= a[1]+b[5] \\ i=6 \quad a[6] &= a[2]+b[6] \\ i=7 \quad a[7] &= a[3]+b[7] \\ i=8 \quad a[8] &= a[4]+b[8] \\ i=9 \quad a[9] &= a[5]+b[9] \\ i=10 \quad a[10] &= a[6]+b[10] \\ i=11 \quad a[11] &= a[7]+b[11] \end{aligned}$$



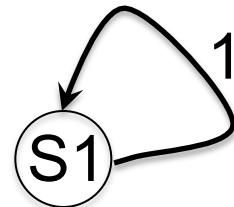
This is also a self-true  
dependence. But ...  
can it be vectorized?



# Cycles in the DG (IV)

```
for (int i=1;i<n;i++){  
S1  a[i]=a[i-1]+b[i];  
}
```

$$\begin{aligned} a[1] &= a[0]+b[1] \\ a[2] &= a[1]+b[2] \\ a[3] &= a[2]+b[3] \end{aligned}$$

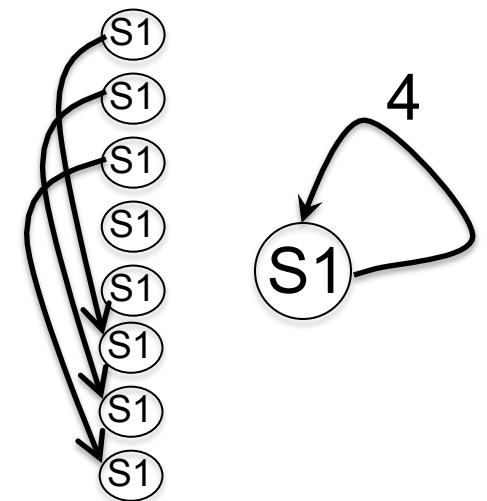


Self true-dependence  
**cannot** be vectorized



```
for (int i=4;i<LEN;i++){  
  a[i]=a[i-4]+b[i];  
}  
  
i=4 a[4] =a[0]+b[4]  
i=5 a[5] =a[1]+b[5]  
i=6 a[6] =a[2]+b[6]  
i=7 a[7] =a[3]+b[7]  
i=8 a[8] =a[4]+b[8]  
i=9 a[9] =a[5]+b[9]  
i=10 a[10]=a[6]+b[10]  
i=11 a[11]=a[7]+b[11]
```

Yes, it can be vectorized because the dependence distance is 4, which is the number of iterations that the SIMD unit can execute simultaneously.



# Cycles in the DG (IV)

S119

```
for (int i=4;i<LEN;i++){
    a[i]=a[i-4]+b[i];
}
```

## Intel Nehalem

**Compiler report:** Loop was vectorized

**Exec. Time scalar code:** 8.4

**Exec. Time vector code:** 3.9

**Speedup:** 2.1

## IBM Power 7

**Compiler report:** Loop was SIMD vectorized

**Exec. Time scalar code:** 6.6

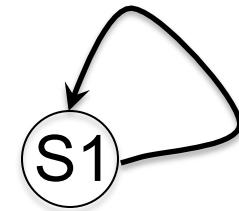
**Exec. Time vector code:** 1.8

**Speedup:** 3.7



# Cycles in the DG (V)

```
for (int i = 0; i < LEN-1; i++) {  
    for (int j = 0; j < LEN; j++)  
S1    a[i+1][j] = a[i][j] + b;  
}
```



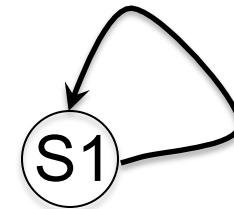
Can this loop be vectorized?

```
i=0, j=0: a[1][0] = a[0][0] + b  
          j=1: a[1][1] = a[0][1] + b  
          j=2: a[1][2] = a[0][2] + b  
i=1  j=0: a[2][0] = a[1][0] + b  
          j=1: a[2][1] = a[1][1] + b  
          j=2: a[2][2] = a[1][2] + b
```



# Cycles in the DG (V)

```
for (int i = 0; i < LEN-1; i++) {  
    for (int j = 0; j < LEN; j++)  
S1    a[i+1][j] = a[i][j] + (float) 1.0;  
}
```



Can this loop be vectorized?

i=0, j=0:  $a[1][0] = a[0][0] + 1$   
j=1:  $a[1][1] = a[0][1] + 1$   
j=2:  $a[1][2] = a[0][2] + 1$   
i=1 j=0:  $a[2][0] = a[1][0] + 1$   
j=1:  $a[2][1] = a[1][1] + 1$   
j=2:  $a[2][2] = a[1][2] + 1$

Dependences occur in the outermost loop.  
- outer loop runs serially  
- inner loop can be vectorized

```
for (int i=0;i<LEN;i++){  
    a[i+1][0:LEN-1]=a[i][0:LEN-1]+b;  
}
```



# Cycles in the DG (V)

S121

```
for (int i = 0; i < LEN-1; i++) {  
    for (int j = 0; j < LEN; j++)  
        a[i+1][j] = a[i][j] + 1;  
}
```

## Intel Nehalem

**Compiler report:** Loop was  
vectorized

**Exec. Time scalar code:** 11.6

**Exec. Time vector code:** 3.2

**Speedup:** 3.5

## IBM Power 7

**Compiler report:** Loop was SIMD  
vectorized

**Exec. Time scalar code:** 3.9

**Exec. Time vector code:** 1.8

**Speedup:** 2.1

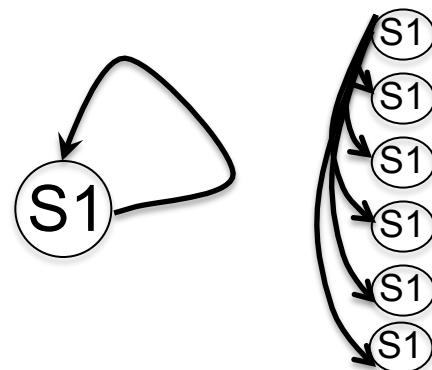


# Cycles in the DG (VI)

- Cycles can appear because the compiler does not know if there are dependences

```
for (int i=0;i<LEN;i++){  
    S1  a[r[i]] = a[r[i]] * (float) 2.0;  
}
```

Is there a value of  $i$  such  
that  $r[i'] = r[i]$ , such that  $i' \neq i$ ?



Compiler cannot resolve the system

To be safe, it considers that a data dependence is possible for every instance of S1

# Cycles in the DG (VI)

- The compiler is conservative.
- The compiler only vectorizes when it can prove that it is safe to do it.

```
for (int i=0;i<LEN;i++){  
    r[i] = i;  
    a[r[i]] = a[r[i]]* (float) 2.0;  
}
```

Does the compiler use the info that  $r[i] = i$  to compute data dependences?



# Cycles in the DG (VI)

S122

```
for (int i=0;i<LEN;i++){  
    a[r[i]]=a[r[i]]*(float)2.0;  
}
```

S123

```
for (int i=0;i<LEN;i++){  
    r[i] = i;  
    a[r[i]]=a[r[i]]*(float)2.0;  
}
```

Does the compiler uses the info that  $r[i] = i$  to compute data dependences?

S122

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Existence of vector dependence  
**Exec. Time scalar code:** 5.0  
**Exec. Time vector code:** --  
**Speedup:** --

S123

**Intel Nehalem**  
**Compiler report:** Partial Loop was vectorized  
**Exec. Time scalar code:** 5.8  
**Exec. Time vector code:** 5.7  
**Speedup:** 1.01



# Cycles in the DG (VI)

S122

```
for (int i=0;i<LEN;i++){  
    a[r[i]]=a[r[i]]*(float)2.0;  
}
```

S123

```
for (int i=0;i<LEN;i++){  
    r[i] = i;  
    a[r[i]]=a[r[i]]*(float)2.0;  
}
```

Does the compiler uses the info that  $r[i] = i$  to compute data dependences?

S122

**IBM Power 7**  
**Compiler report:** Loop was not vectorized because a data dependence prevents SIMD vectorization  
**Exec. Time scalar code:** 2.6  
**Exec. Time vector code:** 2.3  
**Speedup:** 1.1

S123

**IBM Power 7**  
**Compiler report:** Loop was SIMD vectorized  
**Exec. Time scalar code:** 2.1  
**Exec. Time vector code:** 0.9  
**Speedup:** 2.3



# Dependence Graphs and Compiler Vectorization

- No dependences: Vectorized by the compiler
- Acyclic graphs:
  - All dependences are forward:
    - Vectorized by the compiler
  - Some backward dependences:
    - Sometimes vectorized by the compiler
- Cycles in the dependence graph
  - Self-antidependence:
    - Vectorized by the compiler
  - Recurrence:
    - Usually not vectorized by the compiler
  - Other examples



# Loop Transformations

- Compiler Directives
- Loop Distribution or loop fission
- Reordering Statements
- Node Splitting
- Scalar expansion
- Loop Peeling
- Loop Fusion
- Loop Unrolling
- Loop Interchanging



# Compiler Directives (I)

- When the compiler does not vectorize automatically due to dependences the programmer can inform the compiler that it is safe to vectorize:

```
#pragma ivdep (ICC compiler)
```

```
#pragma ibm independent_loop (XLC compiler)
```

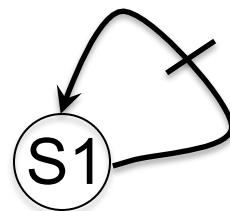


# Compiler Directives (I)

- This loop can be vectorized when  $k < -3$  and  $k \geq 0$ .
- Programmer knows that  $k \geq 0$

```
for (int i=val;i<LEN-k;i++)  
    a[i]=a[i+k]+b[i];
```

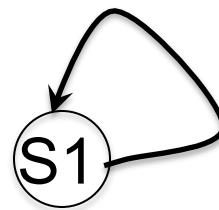
If ( $k \geq 0$ ) → no dependence or self-anti-dependence



$k = 1$

$$\begin{aligned}a[0] &= a[1] + b[0] \\a[1] &= a[2] + b[1] \\a[2] &= a[3] + b[2]\end{aligned}$$

Can  
be vectorized



$k = -1$

$$\begin{aligned}a[1] &= a[0] + b[0] \\a[2] &= a[1] + b[1] \\a[3] &= a[2] + b[2]\end{aligned}$$

Cannot  
be vectorized

If ( $k < 0$ ) → self-true dependence



# Compiler Directives (I)

- This loop can be vectorized when  $k < -3$  and  $k \geq 0$ .
- Programmer knows that  $k \geq 0$

How can the programmer tell the compiler that  $k \geq 0$

```
for (int i=val;i<LEN-k;i++)  
    a[i]=a[i+k]+b[i];
```



# Compiler Directives (I)

- This loop can be vectorized when  $k < -3$  and  $k \geq 0$ .
- Programmer knows that  $k \geq 0$

Intel ICC provides the `#pragma ivdep` to tell the compiler that it is safe to ignore unknown dependences

```
#pragma ivdep
for (int i=val;i<LEN-k;i++)
    a[i]=a[i+k]+b[i];
```

wrong results will be obtained if loop is vectorized when  $-3 < k < 0$



# Compiler Directives (I)

S124

```
for (int i=0;i<LEN-k;i++)  
    a[i]=a[i+k]+b[i];
```

S124\_1

```
if (k>=0)  
    for (int i=0;i<LEN-k;i++)  
        a[i]=a[i+k]+b[i];  
if (k<0)  
    for (int i=0;i<LEN-k;i++)  
        a[i]=a[i+k]+b[i];
```

S124\_2

```
if (k>=0)  
#pragma ivdep  
for (int i=0;i<LEN-k;i++)  
    a[i]=a[i+k]+b[i];  
if (k<0)  
    for (int i=0;i<LEN-k;i++)  
        a[i]=a[i+k]+b[i];
```

S124 and S124\_1

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Existence of vector dependence  
**Exec. Time scalar code:** 6.0  
**Exec. Time vector code:** --  
**Speedup:** --

S124\_2

**Intel Nehalem**  
**Compiler report:** Loop was vectorized  
**Exec. Time scalar code:** 6.0  
**Exec. Time vector code:** 2.4  
**Speedup:** 2.5



# Compiler Directives (I)

S124

```
for (int i=0;i<LEN-k;i++) if (k>=0)
    a[i]=a[i+k]+b[i];
    for (int i=0;i<LEN-k;i++)
        a[i]=a[i+k]+b[i];
if (k<0)
    for (int i=0);i<LEN-k;i++) if (k<0)
        a[i]=a[i+k]+b[i];
```

S124\_1

```
for (int i=0;i<LEN-k;i++)
    a[i]=a[i+k]+b[i];
if (k<0)
    for (int i=0);i<LEN-k;i++) if (k<0)
        a[i]=a[i+k]+b[i];
```

S124\_2

```
if (k>=0)
#pragma ibm independent_loop
for (int i=0;i<LEN-k;i++)
    a[i]=a[i+k]+b[i];
for (int i=0);i<LEN-k;i++) if (k<0)
    a[i]=a[i+k]+b[i];
```

S124 and S124\_1

**IBM Power 7**

**Compiler report:** Loop was not vectoriced because a data dependence prevents SIMD vectorization

**Exec. Time scalar code:** 2.2

**Exec. Time vector code:** --

**Speedup:** --

S124\_2

```
#pragma ibm independent_loop
needs AIX OS (we ran the
experiments on Linux)
```



# Compiler Directives (II)

- Programmer can disable vectorization of a loops when the when the vector code runs slower than the scalar code

```
#pragma novector (ICC compiler)
```

```
#pragma nosimd (XLC compiler)
```

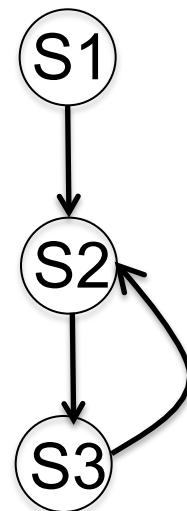


# Compiler Directives (II)

Vector code can run slower than scalar code

```
for (int i=1;i<LEN;i++){  
S1  a[i] = b[i] + c[i];  
S2  d[i] = a[i] + e[i-1];  
S3  e[i] = d[i] + c[i];  
}
```

Less locality when  
executing in vector mode



S1 can be vectorized  
S2 and S3 cannot be vectorized (as they are)



# Compiler Directives (II)

S116

```
#pragma novector  
for (int i=1;i<LEN;i++){  
    a[i] = b[i] + c[i];  
    d[i] = a[i] + e[i-1];  
    e[i] = d[i] + c[i];  
}
```

S116

**Intel Nehalem**

**Compiler report:** Loop was partially vectorized

**Exec. Time scalar code:** 14.7

**Exec. Time vector code:** 18.1

**Speedup:** 0.8



# Loop Distribution

- It is also called loop fission.
- Divides loop control over different statements in the loop body.

```
for (i=1; i<LEN; i++) {           for (i=1; i<LEN; i++)  
    a[i]= (float)sqrt(b[i]) +      a[i]= (float)sqrt(b[i]) +  
        (float)sqrt(c[i]);          (float)sqrt(c[i]);  
    dummy(a,b,c);                →  
}  
                                for (i=1; i<LEN; i++)  
                                dummy(a,b,c);
```

- Compiler cannot analyze the dummy function.  
As a result, the compiler cannot apply loop distribution,  
because it does not know if it is a legal transformation
- Programmer can apply loop distribution if legal.



# Loop Distribution

S126

```
for (i=1; i<LEN; i++) {  
    a[i]= (float)sqrt(b[i]) +  
           (float)sqrt(c[i]);  
    dummy(a,b,c);  
}
```

S126\_1

```
for (i=1; i<LEN; i++)  
    a[i]= (float)sqrt(b[i]) +  
           (float)sqrt(c[i]);  
for (i=1; i<LEN; i++)  
    dummy(a,b,c);
```

S126

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized  
**Exec. Time scalar code:** 4.3  
**Exec. Time vector code:** --  
**Speedup:** --

S126\_1

**Intel Nehalem**  
**Compiler report:**  
- Loop 1 was vectorized.  
- Loop 2 was not vectorized  
**Exec. Time scalar code:** 5.1  
**Exec. Time vector code:** 1.1  
**Speedup:** 4.6



# Loop Distribution

S126

```
for (i=1; i<LEN; i++) {  
    a[i]= (float)sqrt(b[i]) +  
           (float)sqrt(c[i]);  
    dummy(a,b,c);  
}
```

S126\_1

```
for (i=1; i<LEN; i++)  
    a[i]= (float)sqrt(b[i]) +  
           (float)sqrt(c[i]);  
for (i=1; i<LEN; i++)  
    dummy(a,b,c);
```

S126

**IBM Power 7**  
**Compiler report:** Loop was not SIMD vectorized  
**Exec. Time scalar code:** 1.3  
**Exec. Time vector code:** --  
**Speedup:** --

S126\_1

**IBM Power 7**  
**Compiler report:**

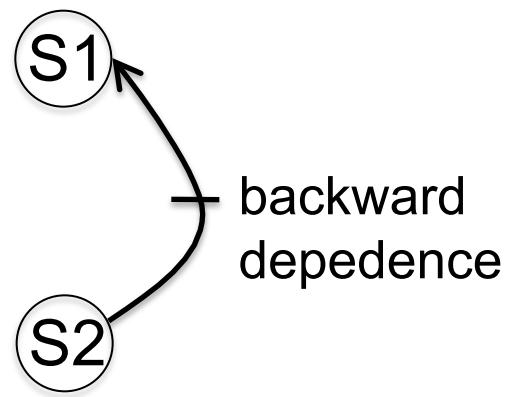
- Loop 1 was SIMD vectorized.
- Loop 2 was not SIMD vectorized

**Exec. Time scalar code:** 1.14  
**Exec. Time vector code:** 1.0  
**Speedup:** 1.14

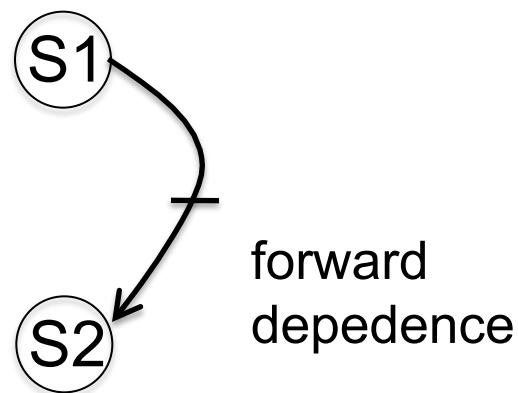


# Reordering Statements

```
for (i=0; i<LEN; i++) {  
S1  a[i]= b[i] + c[i];  
S2  d[i] = a[i+1]+(float)1.0;  
}
```



```
for (i=0; i<LEN; i++) {  
S1  d[i] = a[i+1]+(float)1.0;  
S2  a[i]= b[i] + c[i];  
}
```



# Reordering Statements

S114

```
for (i=0; i<LEN; i++) {  
    a[i]= b[i] + c[i];  
    d[i] = a[i+1]+(float)1.0;  
}
```

S114\_1

```
for (i=0; i<LEN; i++) {  
    d[i] = a[i+1]+(float)1.0;  
    a[i]= b[i] + c[i];  
}
```

S114

S114\_1

## Intel Nehalem

**Compiler report:** Loop was not vectorized. Existence of vector dependence

**Exec. Time scalar code:** 12.6

**Exec. Time vector code:** --

**Speedup:** --

## Intel Nehalem

**Compiler report:** Loop was vectorized.

**Exec. Time scalar code:** 10.7

**Exec. Time vector code:** 9.4

**Speedup:** 1.03



# Reordering Statements

S114

```
for (i=0; i<LEN; i++) {  
    a[i]= b[i] + c[i];  
    d[i] = a[i+1]+(float)1.0;  
}
```

S114\_1

```
for (i=0; i<LEN; i++) {  
    d[i] = a[i+1]+(float)1.0;  
    a[i]= b[i] + c[i];  
}
```

The IBM XLC compiler generated the same code in both cases

S114

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 3.3  
**Exec. Time vector code:** 1.8  
**Speedup:** 1.8

S114\_1

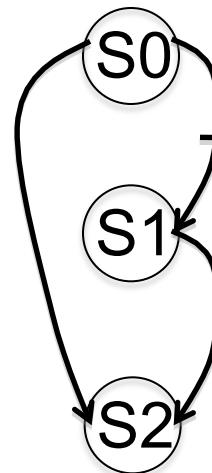
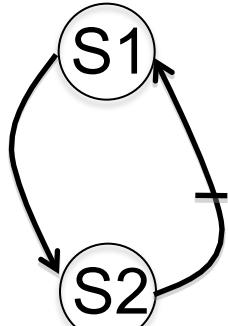
**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 3.3  
**Exec. Time vector code:** 1.8  
**Speedup:** 1.8



# Node Splitting

```
for (int i=0;i<LEN-1;i++){  
S1 a[i]=b[i]+c[i];  
S2 d[i]=(a[i]+a[i+1])*(float)0.5;  
}
```

```
for (int i=0;i<LEN-1;i++){  
S0 temp[i]=a[i+1];  
S1 a[i]=b[i]+c[i];  
S2 d[i]=(a[i]+temp[i])*(float) 0.5  
}
```



# Node Splitting

S126

```
for (int i=0;i<LEN-1;i++){
    a[i]=b[i]+c[i];
    d[i]=(a[i]+a[i+1])*(float)0.5;
}
```

S126\_1

```
for (int i=0;i<LEN-1;i++){
    temp[i]=a[i+1];
    a[i]=b[i]+c[i];
    d[i]=(a[i]+temp[i])*(float)0.5;
}
```

S126

## Intel Nehalem

**Compiler report:** Loop was not vectorized. Existence of vector dependence

**Exec. Time scalar code:** 12.6

**Exec. Time vector code:** --

**Speedup:** --

S126\_1

## Intel Nehalem

**Compiler report:** Loop was vectorized.

**Exec. Time scalar code:** 13.2

**Exec. Time vector code:** 9.7

**Speedup:** 1.3



# Node Splitting

S126

```
for (int i=0;i<LEN-1;i++){  
S1 a[i]=b[i]+c[i];  
S2 d[i]=(a[i]+a[i+1])*(float)0.5;  
}
```

S126\_1

```
for (int i=0;i<LEN-1;i++){  
S0 temp[i]=a[i+1];  
S1 a[i]=b[i]+c[i];  
S2 d[i]=(a[i]+temp[i])*(float) 0.5  
}
```

S126

**IBM Power 7**  
**Compiler report:** Loop was SIMD vectorized  
**Exec. Time scalar code:** 3.8  
**Exec. Time vector code:** 1.7  
**Speedup:** 2.2

S126\_1

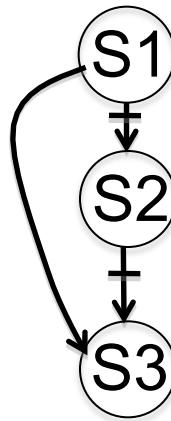
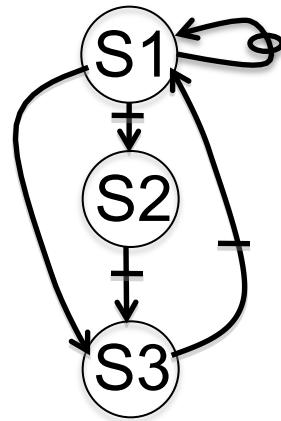
**IBM Power 7**  
**Compiler report:** Loop was SIMD vectorized  
**Exec. Time scalar code:** 5.1  
**Exec. Time vector code:** 2.4  
**Speedup:** 2.0



# Scalar Expansion

```
for (int i=0;i<n;i++){  
S1 t = a[i];  
S2 a[i] = b[i];  
S3 b[i] = t;  
}
```

```
for (int i=0;i<n;i++){  
S1 t[i] = a[i];  
S2 a[i] = b[i];  
S3 b[i] = t[i];  
}
```



# Scalar Expansion

S139

```
for (int i=0;i<n;i++){  
    t = a[i];  
    a[i] = b[i];  
    b[i] = t;  
}
```

S139\_1

```
for (int i=0;i<n;i++){  
    t[i] = a[i];  
    a[i] = b[i];  
    b[i] = t[i];  
}
```

S139

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 0.7  
**Exec. Time vector code:** 0.4  
**Speedup:** 1.5

S139\_1

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 0.7  
**Exec. Time vector code:** 0.4  
**Speedup:** 1.5



# Scalar Expansion

S139

```
for (int i=0;i<n;i++){  
    t = a[i];  
    a[i] = b[i];  
    b[i] = t;  
}
```

S139\_1

```
for (int i=0;i<n;i++){  
    t[i] = a[i];  
    a[i] = b[i];  
    b[i] = t[i];  
}
```

S139

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 0.28  
**Exec. Time vector code:** 0.14  
**Speedup:** 2

S139\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 0.28  
**Exec. Time vector code:** 0.14  
**Speedup:** 2.0



# Loop Peeling

- Remove the first/s or the last/s iteration of the loop into separate code outside the loop
- It is always legal, provided that no additional iterations are introduced.
- When the trip count of the loop is not constant the peeled loop has to be protected with additional runtime tests.
- This transformation is useful to enforce a particular initial memory alignment on array references prior to loop vectorization.

```
for (i=0; i<LEN; i++) → A[0] = B[0] + C[0];  
    A[i] = B[i] + C[i];           for (i=1; i<LEN; i++)  
                                A[i] = B[i] + C[i];
```



# Loop Peeling

- Remove the first/s or the last/s iteration of the loop into separate code outside the loop
- It is always legal, provided that no additional iterations are introduced.
- When the trip count of the loop is not constant the peeled loop has to be protected with additional runtime tests.
- This transformation is useful to enforce a particular initial memory alignment on array references prior to loop vectorization.

```
if (N>=1)
    A[0] = B[0] + C[0];
for (i=1; i<LEN; i++)
    A[i] = B[i] + C[i];
```

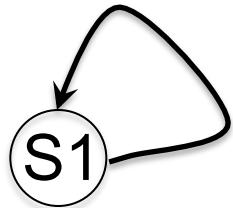
**for (i=0; i<LEN; i++) →**



# Loop Peeling

```
for (int i=0;i<LEN;i++){  
S1  a[i] = a[i] + a[0];  
}
```

$$\begin{aligned}a[0] &= a[0] + a[0] \\a[1] &= a[1] + a[0] \\a[2] &= a[2] + a[0]\end{aligned}$$



Self true-dependence  
**is not vectorized**

```
a[0]= a[0] + a[0];  
for (int i=1;i<LEN;i++){  
  a[i] = a[i] + a[0]  
}
```

After loop peeling, there are no dependences, and the loop can be vectorized

# Loop Peeling

S127

```
for (int i=0;i<LEN;i++){  
S1 a[i] = a[i] + a[0];  
}
```

S127\_1

```
a[0]= a[0] + a[0];  
for (int i=1;i<LEN;i++){  
 a[i] = a[i] + a[0]  
}
```

S127

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Existence of vector dependence  
**Exec. Time scalar code:** 6.7  
**Exec. Time vector code:** --  
**Speedup:** --

S127\_1

**Intel Nehalem**  
**Compiler report:** Loop was vectorized.  
**Exec. Time scalar code:** 6.6  
**Exec. Time vector code:** 1.2  
**Speedup:** 5.2



# Loop Peeling

S127

```
for (int i=0;i<LEN;i++)  
{  
    a[i] = a[i] + a[0];  
}
```

S127\_1

```
a[0]= a[0] + a[0];  
for (int i=1;i<LEN;i++)  
{  
    a[i] = a[i] + a[0];  
}
```

S127\_2

```
a[0]= a[0] + a[0];  
float t = a[0];  
for (int i=1;i<LEN;i++)  
{  
    a[i] = a[i] + t;  
}
```

S127

S127\_1

S127\_2

**IBM Power 7**

**Compiler report:** Loop  
was not SIMD vectorized  
**Time scalar code:** 2.4  
**Time vector code:** --  
**Speedup:** --

**IBM Power 7**

**Compiler report:** Loop  
was not SIMD vectorized  
**Exec. scalar code:** 2.4  
**Exec. vector code:** --  
**Speedup:** --

**IBM Power 7**

**Compiler report:** Loop  
was vectorized  
**Exec. scalar code:** 1.58  
**Exec. vector code:** 0.62  
**Speedup:** 2.54



# Loop Interchanging

- This transformation switches the positions of one loop that is tightly nested within another loop.

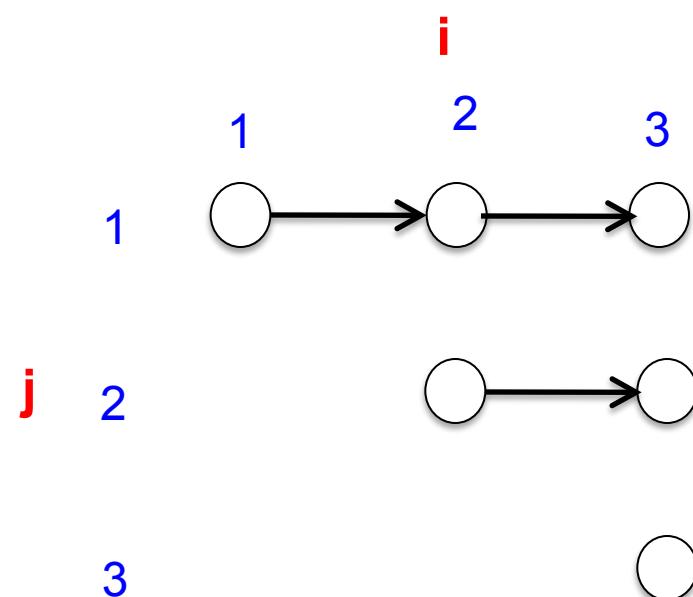
```
for (i=0; i<LEN; i++)
  for (j=0; j<LEN; j++)
    A[i][j]=0.0;
```

```
for (j=0; j<LEN; j++)
  for (i=0; i<LEN; i++)
    A[i][j]=0.0;
```



# Loop Interchanging

```
for (j=1; j<LEN; j++) {
    for (i=j; i<LEN; i++) {
        A[i][j]=A[i-1][j]+(float) 1.0;
    }
}
```



$j=1 \left\{ \begin{array}{l} i=1 A[1][1]=A[0][1] + 1 \\ i=2 A[2][1]=A[1][1] + 1 \\ i=3 A[3][1]=A[2][1] + 1 \end{array} \right.$

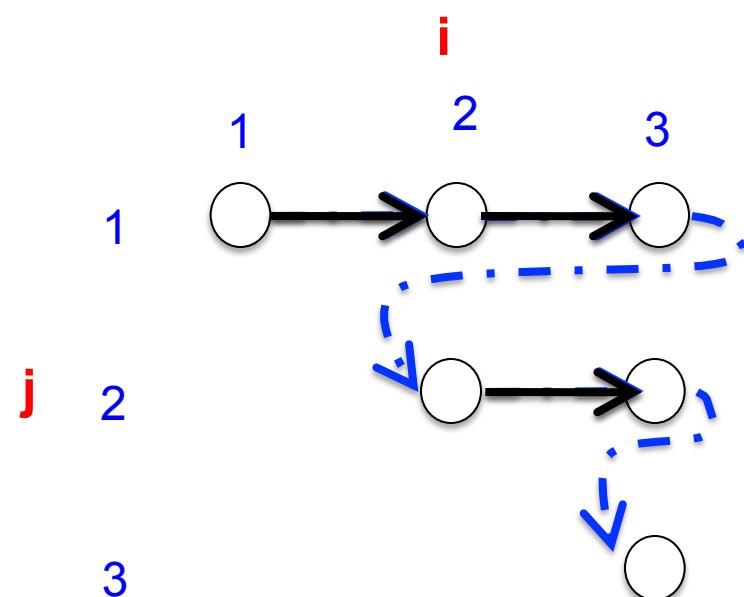
$j=2 \left\{ \begin{array}{l} i=2 A[2][2]=A[1][2] + 1 \\ i=3 A[3][2]=A[2][2] + 1 \end{array} \right.$

$j=3 \quad i=3 A[3][3]=A[2][3] + 1$



# Loop Interchanging

```
for (j=1; j<LEN; j++) {
    for (i=j; i<LEN; i++) {
        A[i][j]=A[i-1][j]+(float) 1.0;
    }
}
```



$j=1 \left\{ \begin{array}{l} i=1 A[1][1]=A[0][1] + 1 \\ i=2 A[2][1]=A[1][1] + 1 \\ i=3 A[3][1]=A[2][1] + 1 \end{array} \right.$

$j=2 \left\{ \begin{array}{l} i=2 A[2][2]=A[1][2] + 1 \\ i=3 A[3][2]=A[2][2] + 1 \end{array} \right.$

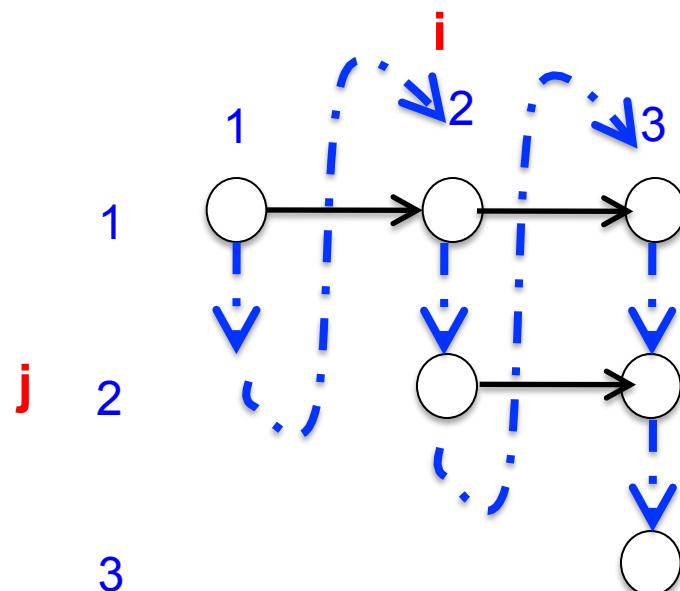
$j=3 \quad i=3 A[3][3]=A[2][3] + 1$

Inner loop cannot be vectorized  
because of self-dependence



# Loop Interchanging

```
for (i=1; i<LEN; i++) {
    for (j=1; j<i+1; j++) {
        A[i][j]=A[i-1][j]+(float) 1.0;
    }
}
```



$i=1 \left\{ \begin{array}{l} j=1 A[1][1]=A[0][1] + 1 \\ \dots \end{array} \right.$

$i=2 \left\{ \begin{array}{l} j=1 A[2][1]=A[1][1] + 1 \\ j=2 A[2][2]=A[1][2] + 1 \\ \dots \end{array} \right.$

$i=3 \left\{ \begin{array}{l} j=1 A[3][1]=A[2][1] + 1 \\ j=2 A[3][2]=A[2][2] + 1 \\ j=3 A[3][3]=A[2][3] + 1 \\ \dots \end{array} \right.$

Loop interchange is legal  
No dependences in inner loop



# Loop Interchanging

S228

```
for (j=1; j<LEN; j++){  
    for (i=j; i<LEN; i++){  
        A[i][j]=A[i-1][j]+(float)1.0;  
    }  
}
```

S228\_1

```
for (i=1; i<LEN; i++){  
    for (j=1; j<i+1; j++){  
        A[i][j]=A[i-1][j]+(float)1.0;  
    }  
}
```

S228

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized.  
**Exec. Time scalar code:** 2.3  
**Exec. Time vector code:** --  
**Speedup:** --

S228\_1

**Intel Nehalem**  
**Compiler report:** Loop was vectorized.  
**Exec. Time scalar code:** 0.6  
**Exec. Time vector code:** 0.2  
**Speedup:** 3



# Loop Interchanging

S228

```
for (j=1; j<LEN; j++){  
    for (i=j; i<LEN; i++){  
        A[i][j]=A[i-1][j]+(float)1.0;  
    }  
}
```

S228\_1

```
for (i=1; i<LEN; i++){  
    for (j=1; j<i+1; j++){  
        A[i][j]=A[i-1][j]+(float)1.0;  
    }  
}
```

S228

**IBM Power 7**  
**Compiler report:** Loop was not SIMD vectorized  
**Exec. Time scalar code:** 0.5  
**Exec. Time vector code:** --  
**Speedup:** --

S228\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD vectorized  
**Exec. Time scalar code:** 0.2  
**Exec. Time vector code:** 0.14  
**Speedup:** 1.42



# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
    - Reductions
  - Data Alignment
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization using intrinsics

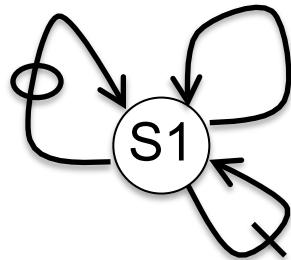


# Reductions

- Reduction is an operation, such as addition, which is applied to the elements of an array to produce a result of a lesser rank.

## Sum Reduction

```
sum =0;  
for (int i=0;i<LEN;++i){  
    sum+= a[i];  
}
```



## Max Loc Reduction

```
x = a[0];  
index = 0;  
for (int i=0;i<LEN;++i){  
    if (a[i] > x) {  
        x = a[i];  
        index = i;  
    }  
}
```



# Reductions

S131

```
sum =0;  
for (int i=0;i<LEN;++i){  
    sum+= a[i];  
}
```

S132

```
x = a[0];  
index = 0;  
for (int i=0;i<LEN;++i){  
    if (a[i] > x) {  
        x = a[i];  
        index = i;  
    }  
}
```

S131

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 5.2  
**Exec. Time vector code:** 1.2  
**Speedup:** 4.1

S132

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 9.6  
**Exec. Time vector code:** 2.4  
**Speedup:** 3.9



# Reductions

S131

```
sum =0;  
for (int i=0;i<LEN;++i){  
    sum+= a[i];  
}
```

S132

```
x = a[0];  
index = 0;  
for (int i=0;i<LEN;++i){  
    if (a[i] > x) {  
        x = a[i];  
        index = i;  
    }  
}
```

S131

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 1.1  
**Exec. Time vector code:** 0.4  
**Speedup:** 2.4

S132

**IBM Power 7**  
**Compiler report:** Loop was not  
SIMD vectorized  
**Exec. Time scalar code:** 4.4  
**Exec. Time vector code:** --  
**Speedup:** --



# Reductions

S141\_1

```
for (int i = 0; i < 64; i++){
    max[i] = a[i];
    loc[i] = i; }
for (int i = 0; i < LEN; i+=64){
    for (int j=0, k=i; k<i+64; k++,j++)
{
    int cmp = max[j] < a[k];
    max[j] = cmp ? a[k] : max[j];
    loc[j] = cmp ? k : loc[j];
} }
MAX = max[0];
LOC = 0;
for (int i = 0; i < 64; i++){
    if (MAX < max[i]){
        MAX = max[i];
        LOC = loc[i];
    } }
```

S141\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 10.2  
**Exec. Time vector code:** 2.7  
**Speedup:** 3.7

S141\_2

**IBM Power 7**  
**A version written with intrinsics**  
**runs in 1.6 secs.**



# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
    - Induction variables
  - Data Alignment
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics

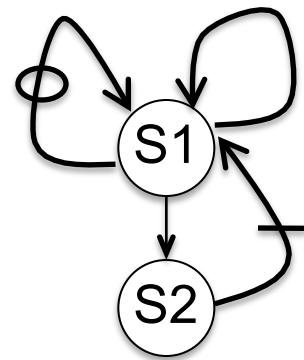


# Induction variables

- Induction variable is a variable that can be expressed as a function of the loop iteration variable

```
float s = (float)0.0;  
for (int i=0;i<LEN;i++){  
    s += (float)2.;  
    a[i] = s * b[i];  
}
```

```
for (int i=0;i<LEN;i++){  
    a[i] = (float)2.*(i+1)*b[i];  
}
```



# Induction variables

S133

```
float s = (float)0.0;  
for (int i=0;i<LEN;i++){  
    s += (float)2.;  
    a[i] = s * b[i];  
}
```

S133\_1

```
for (int i=0;i<LEN;i++){  
    a[i] = (float)2.*(i+1)*b[i];  
}
```

The Intel ICC compiler generated the same vector code in both cases

S133

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 6.1  
**Exec. Time vector code:** 1.9  
**Speedup:** 3.1

S133\_1

**Intel Nehalem**  
**Compiler report:** Loop was  
vectorized.  
**Exec. Time scalar code:** 8.4  
**Exec. Time vector code:** 1.9  
**Speedup:** 4.2



# Induction variables

S133

```
float s = (float)0.0;  
for (int i=0;i<LEN;i++){  
    s += (float)2.;  
    a[i] = s * b[i];  
}
```

S133\_1

```
for (int i=0;i<LEN;i++){  
    a[i] = (float)2.*(i+1)*b[i];  
}
```

S133

**IBM Power 7**  
**Compiler report:** Loop was not SIMD vectorized  
**Exec. Time scalar code:** 2.7  
**Exec. Time vector code:** --  
**Speedup:** --

S133\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD vectorized  
**Exec. Time scalar code:** 3.7  
**Exec. Time vector code:** 1.4  
**Speedup:** 2.6



# Induction Variables

- Coding style matters:

```
for (int i=0;i<LEN;i++) {      for (int i=0;i<LEN;i++){  
    *a = *b + *c;  
    a++; b++; c++;  
}
```

These codes are equivalent, but ...



# Induction Variables

S134

```
for (int i=0;i<LEN;i++) {  
    *a = *b + *c;  
    a++; b++; c++;  
}
```

S134\_1

```
for (int i=0;i<LEN;i++){  
    a[i] = b[i] + c[i];  
}
```

S134

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized.  
**Exec. Time scalar code:** 5.5  
**Exec. Time vector code:** --  
**Speedup:** --

S134\_1

**Intel Nehalem**  
**Compiler report:** Loop was vectorized.  
**Exec. Time scalar code:** 6.1  
**Exec. Time vector code:** 3.2  
**Speedup:** 1.8



# Induction Variables

S134

```
for (int i=0;i<LEN;i++) {  
    *a = *b + *c;  
    a++; b++; c++;  
}
```

S134\_1

```
for (int i=0;i<LEN;i++){  
    a[i] = b[i] + c[i];  
}
```

The IBM XLC compiler generated the same code in both cases

S134

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 2.2  
**Exec. Time vector code:** 1.0  
**Speedup:** 2.2

S134\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 2.2  
**Exec. Time vector code:** 1.0  
**Speedup:** 2.2



# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
  - **Data Alignment**
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics

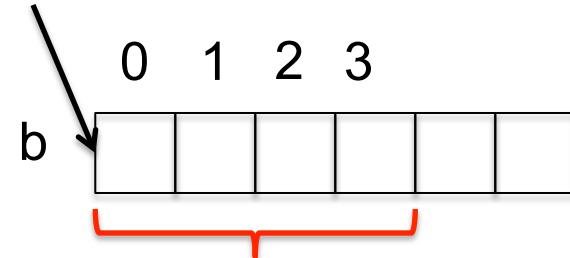


# Data Alignment

- Vector loads/stores load/store 128 consecutive bits to a vector register.
- Data addresses need to be 16-byte (128 bits) aligned to be loaded/stored
  - Intel platforms support aligned and unaligned load/stores
  - IBM platforms do not support unaligned load/stores

```
void test1(float *a, float *b, float *c)
{
    for (int i=0; i<LEN; i++){
        a[i] = b[i] + c[i];
    }
}
```

Is &b[0] 16-byte aligned?



vector load loads b[0] ... b[3]



# Data Alignment

- To know if a pointer is 16-byte aligned, the last digit of the pointer address in hex must be 0.
- Note that if `&b[0]` is 16-byte aligned, and is a single precision array, then `&b[4]` is also 16-byte aligned

```
__attribute__ ((aligned(16))) float b[1024];
```

```
int main(){
    printf("%p, %p\n", &b[0], &b[4]);
}
```

Output:

0x7fff1e9d8580, 0x7fff1e9d8590



# Data Alignment

- In many cases, the compiler cannot statically know the alignment of the address in a pointer
- The compiler assumes that the base address of the pointer is 16-byte aligned and adds a run-time checks for it
  - if the runtime check is false, then it uses another code (which may be scalar)



# Data Alignment

- Manual 16-byte alignment can be achieved by forcing the base address to be a multiple of 16.

```
__attribute__((aligned(16))) float b[N];
float* a = (float*) memalign(16,N*sizeof(float));
```

- When the pointer is passed to a function, the compiler should be aware of where the 16-byte aligned address of the array starts.

```
void func1(float *a, float *b,
float *c) {
    __assume_aligned(a, 16);
    __assume_aligned(b, 16);
    __assume_aligned(c, 16);
    for int (i=0; i<LEN; i++) {
        a[i] = b[i] + c[i];
    }
}
```



# Data Alignment - Example

```
float A[N] __attribute__((aligned(16)));
float B[N] __attribute__((aligned(16)));
float C[N] __attribute__((aligned(16)));

void test(){
    for (int i = 0; i < N; i++){
        C[i] = A[i] + B[i];
    }
}
```



# Data Alignment - Example

```
float A[N] __attribute__((aligned(16)));
float B[N] __attribute__((aligned(16)));
float C[N] __attribute__((aligned(16)));

void test1(){
    __m128 rA, rB, rC;
    for (int i = 0; i < N; i+=4){
        rA = _mm_load_ps(&A[i]);
        rB = _mm_load_ps(&B[i]);
        rC = _mm_add_ps(rA, rB);
        _mm_store_ps(&C[i], rC);
    }
}

void test3(){
    __m128 rA, rB, rC;
    for (int i = 1; i < N-3; i+=4){
        rA = _mm_loadu_ps(&A[i]);
        rB = _mm_loadu_ps(&B[i]);
        rC = _mm_add_ps(rA, rB);
        _mm_storeu_ps(&C[i], rC);
    }
}
```

```
void test2(){
    __m128 rA, rB, rC;
    for (int i = 0; i < N; i+=4){
        rA = _mm_loadu_ps(&A[i]);
        rB = _mm_loadu_ps(&B[i]);
        rC = _mm_add_ps(rA, rB);
        _mm_storeu_ps(&C[i], rC);
    }
}
```

Nanosecond per iteration			
	Core 2 Duo	Intel i7	Power 7
Aligned	0.577	0.580	0.156
Aligned (unaligned Id)	0.689	0.581	0.241
Unaligned	2.176	0.629	0.243



# Alignment in a struct

```
struct st{  
    char A;  
    int B[64];  
    float C;  
    int D[64];  
};  
  
int main(){  
    st s1;  
    printf("%p, %p, %p, %p\n", &s1.A, s1.B, &s1.C, s1.D);}
```

Output:

0x7ffe6765f00, 0x7ffe6765f04, 0x7ffe6766004, 0x7ffe6766008

- Arrays B and D are not 16-bytes aligned (see the address)



# Alignment in a struct

```
struct st{  
    char A;  
    int B[64] __attribute__((aligned(16)));  
    float C;  
    int D[64] __attribute__((aligned(16)));  
};  
  
int main(){  
    st s1;  
    printf("%p, %p, %p, %p\n", &s1.A, s1.B, &s1.C, s1.D);}
```

Output:

0x7fff1e9d8580, 0x7fff1e9d8590, 0x7fff1e9d8690, 0x7fff1e9d86a0

- Arrays A and B are aligned to 16-bytes (notice the 0 in the 4 least significant bits of the address)
- Compiler automatically does padding



# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
  - Data Alignment
  - **Aliasing**
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics



# Aliasing

- Can the compiler vectorize this loop?

```
void func1(float *a, float *b, float *c){  
    for (int i = 0; i < LEN; i++) {  
        a[i] = b[i] + c[i];  
    }  
}
```



# Aliasing

- Can the compiler vectorize this loop?

```
float* a = &b[1];  
...  
void func1(float *a, float *b, float *c)  
{  
    for (int i = 0; i < LEN; i++)  
        a[i] = b[i] + c[i];  
}
```

$$\begin{aligned}b[1] &= b[0] + c[0] \\b[2] &= b[1] + c[1]\end{aligned}$$



# Aliasing

- Can the compiler vectorize this loop?

```
float* a = &b[1];  
...  
void func1(float *a, float *b, float *c)  
{  
    for (int i = 0; i < LEN; i++)  
        a[i] = b[i] + c[i];  
}  
  
a and b are aliasing  
There is a self-true dependence  
Vectorizing this loop would  
be illegal
```



# Aliasing

- To vectorize, the compiler needs to guarantee that the pointers are not aliased.
- When the compiler does not know if two pointer are alias, it still vectorizes, but needs to add up-to  $O(n^{12})$  run-time checks, where  $n$  is the number of pointers

When the number of pointers is large, the compiler may decide to not vectorize

```
void func1(float *a, float *b, float *c){  
    for (int i=0; i<LEN; i++)  
        a[i] = b[i] + c[i];  
}
```



# Aliasing

- Two solutions can be used to avoid the run-time checks
  1. static and global arrays
  2. `__restrict__` attribute



# Aliasing

## 1. Static and Global arrays

```
__attribute__((aligned(16))) float a[LEN];
__attribute__((aligned(16))) float b[LEN];
__attribute__((aligned(16))) float c[LEN];

void func1(){
    for (int i=0; i<LEN; i++)
        a[i] = b[i] + c[i];
}

int main() {
    ...
    func1();
}
```



# Aliasing

## 1. `__restrict__` keyword

```
void func1(float* __restrict__ a, float* __restrict__ b,
float* __restrict__ c) {
    __assume_aligned(a, 16);
    __assume_aligned(b, 16);
    __assume_aligned(c, 16);
    for int (i=0; i<LEN; i++)
        a[i] = b[i] + c[i];
}
int main() {
    float* a=(float*) memalign(16,LEN*sizeof(float));
    float* b=(float*) memalign(16,LEN*sizeof(float));
    float* c=(float*) memalign(16,LEN*sizeof(float));
    ...
    func1(a,b,c);
}
```



# Aliasing – Multidimensional arrays

- Example with 2D arrays: pointer-to-pointer declaration.

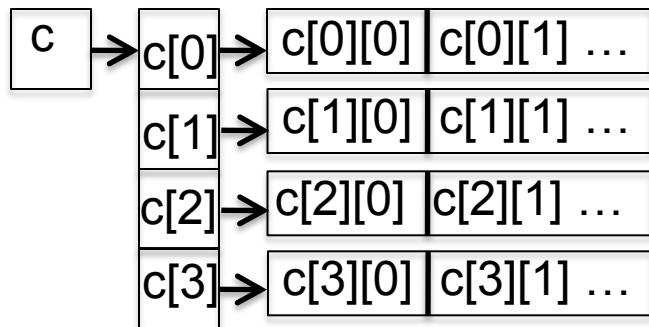
```
void func1(float** __restrict__ a, float**  
__restrict__ b, float** __restrict__ c) {  
    for (int i=0; i<LEN; i++)  
        for (int j=1; j<LEN; j++)  
            a[i][j] = b[i][j-1] * c[i][j];  
}
```



# Aliasing – Multidimensional arrays

- Example with 2D arrays: pointer-to-pointer declaration.

```
void func1(float** __restrict__ a, float** __restrict__  
b, float** __restrict__ c) {  
    for (int i=0; i<LEN; i++)  
        for (int j=1; j<LEN; j++)  
            a[i][j] = b[i][j-1] * c[i][j];  
}
```



`__restrict__` only qualifies  
the first dereferencing of `c`;

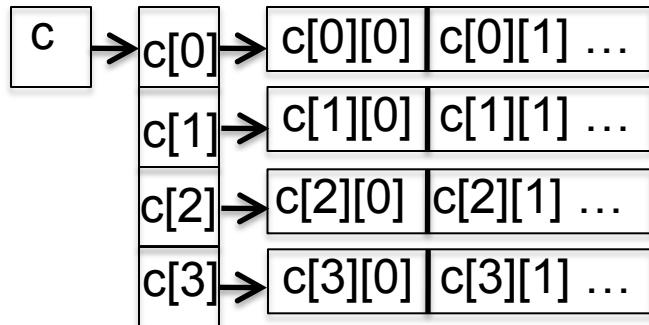
Nothing is said about the  
arrays that can be accessed  
through `c[i]`



# Aliasing – Multidimensional arrays

- Example with 2D arrays: pointer-to-pointer declaration.

```
void func1(float** __restrict__ a, float** __restrict__  
b, float** __restrict__ c) {  
    for (int i=0; i<LEN; i++)  
        for (int j=1; j<LEN; j++)  
            a[i][j] = b[i][j-1] * c[i][j];  
}
```



`__restrict__` only qualifies  
the first dereferencing of `c`;

Nothing is said about the  
arrays that can be accessed  
through `c[i]`

Intel ICC compiler, version 11.1 will vectorize this code.

Previous versions of the Intel compiler or compilers from  
other vendors, such as IBM XLC, will not vectorize it.



# Aliasing – Multidemensional Arrays

- Three solutions when `__restrict__` does not enable vectorization
  - 1. Static and global arrays
  - 2. Linearize the arrays and use `__restrict__` keyword
  - 3. Use compiler directives



# Aliasing – Multidimensional arrays

## 1. Static and Global declaration

```
__attribute__((aligned(16))) float a[N][N];
void t(){

    a[i][j]...
}

int main() {

    ...
    t();
}
```



# Aliasing – Multidimensional arrays

## 2. Linearize the arrays

```
void t(float* __restrict__ A){  
    //Access to Element A[i][j] is now A[i*128+j]  
    ...  
}
```

```
int main() {  
    float* A = (float*) memalign(16,128*128*sizeof(float));  
    ...  
    t(A);  
}
```



# Aliasing – Multidimensional arrays

## 3. Use compiler directives:

```
#pragma ivdep (Intel ICC)
#pragma disjoint(IBM XLC)
```

```
void func1(float **a, float **b, float **c) {
    for (int i=0; i<m; i++) {
        #pragma ivdep
        for (int j=0; j<LEN; j++)
            c[i][j] = b[i][j] * a[i][j];
    }
}
```



# Outline

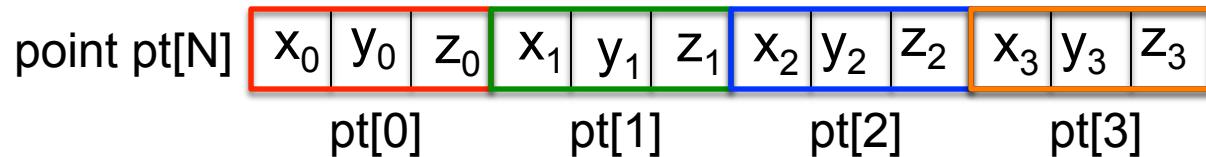
1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
  - Data Alignment
  - Aliasing
  - **Non-unit strides**
  - Conditional Statements
4. Vectorization with intrinsics



# Non-unit Stride – Example I

- Array of a struct

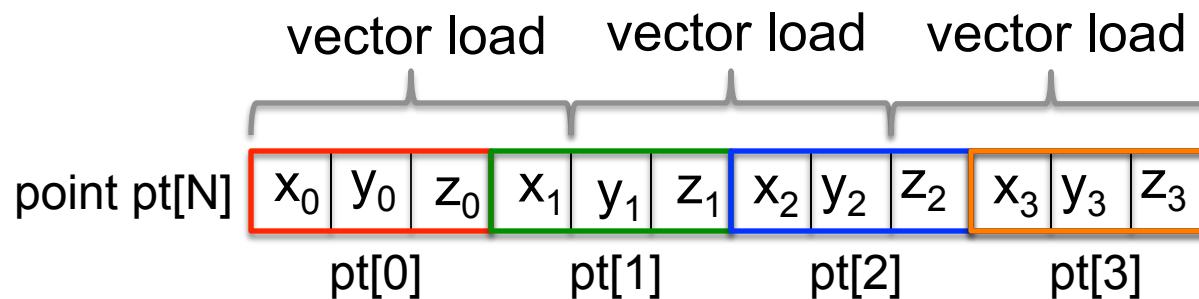
```
typedef struct{int x, y, z}  
point;  
point pt[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pt[i].y *= scale;  
}
```



# Non-unit Stride – Example I

- Array of a struct

```
typedef struct{int x, y, z}  
point;  
point pt[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pt[i].y *= scale;  
}
```

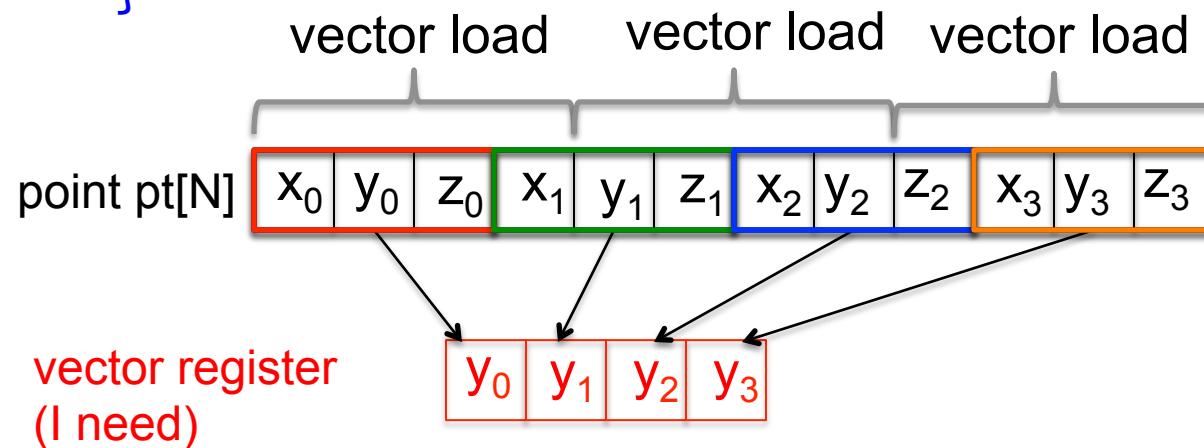


# Non-unit Stride – Example I

- Array of a struct

```
typedef struct{int x, y, z}  
point;  
point pt[LEN];
```

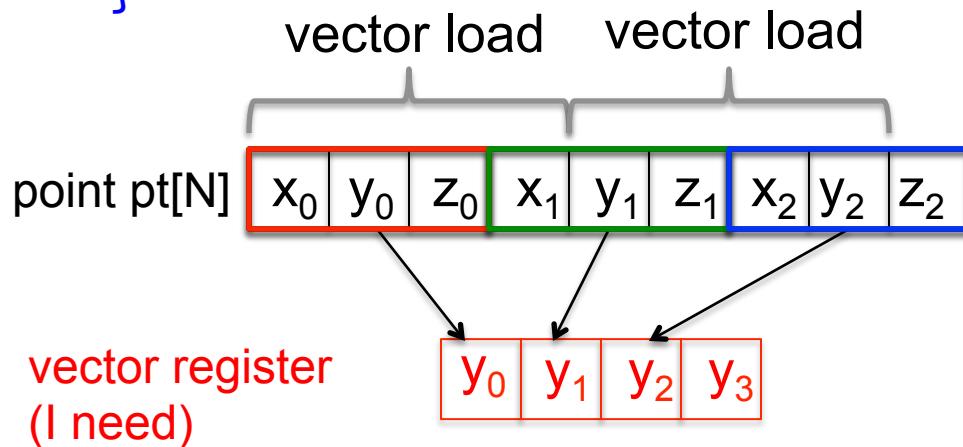
```
for (int i=0; i<LEN; i++) {  
    pt[i].y *= scale;  
}
```



# Non-unit Stride – Example I

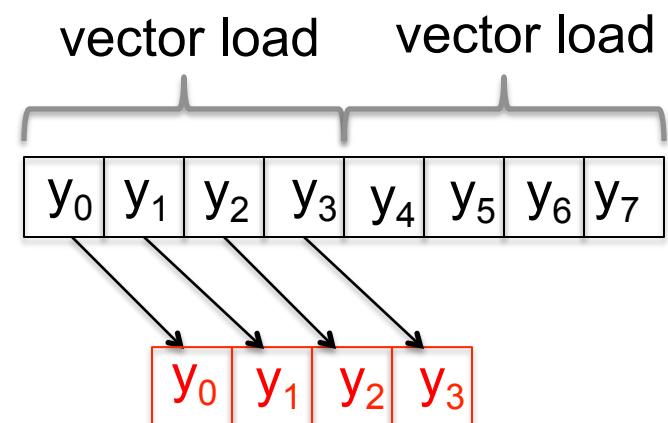
- Array of a struct

```
typedef struct{int x, y, z}  
point;  
point pt[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pt[i].y *= scale;  
}
```



- Arrays

```
int ptx[LEN], int pty[LEN],  
int ptz[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pty[i] *= scale;  
}
```



# Non-unit Stride – Example I

S135

```
typedef struct{int x, y, z}  
point;  
point pt[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pt[i].y *= scale;  
}
```

S135

S135\_1

```
int ptx[LEN], int pty[LEN],  
int ptz[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pty[i] *= scale;  
}
```

S135\_1

## Intel Nehalem

**Compiler report:** Loop was not vectorized. Vectorization possible but seems inefficient

**Exec. Time scalar code:** 6.8  
**Exec. Time vector code:** --  
**Speedup:** --



## Intel Nehalem

**Compiler report:** Loop was vectorized.

**Exec. Time scalar code:** 4.8  
**Exec. Time vector code:** 1.3  
**Speedup:** 3.7

# Non-unit Stride – Example I

S135

```
typedef struct{int x, y, z}  
point;  
point pt[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pt[i].y *= scale;  
}
```

S135

S135\_1

```
int ptx[LEN], int pty[LEN],  
int ptz[LEN];  
  
for (int i=0; i<LEN; i++) {  
    pty[i] *= scale;  
}
```

S135\_1

## IBM Power 7

**Compiler report:** Loop was not SIMD vectorized because it is not profitable to vectorize

**Exec. Time scalar code:** 2.0

**Exec. Time vector code:** --

**Speedup:** --

## IBM Power 7

**Compiler report:** Loop was SIMD vectorized

**Exec. Time scalar code:** 1.8

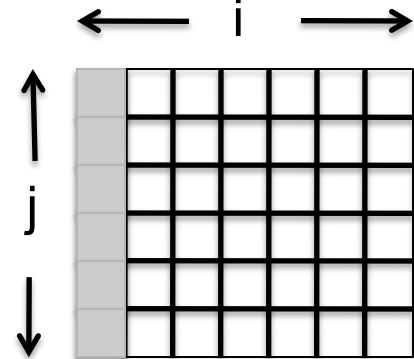
**Exec. Time vector code:** 1.5

**Speedup:** 1.2



# Non-unit Stride – Example II

```
for (int i=0;i<LEN;i++){  
    sum = 0;  
    for (int j=0;j<LEN;j++){  
        sum += A[j][i];  
    }  
    B[i] = sum;  
}
```



```
for (int i=0;i<size;i++){  
    sum[i] = 0;  
    for (int j=0;j<size;j++){  
        sum[i] += A[j][i];  
    }  
    B[i] = sum[i];  
}
```



# Non-unit Stride – Example II

S136

```
for (int i=0;i<LEN;i++){  
    sum = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum += A[j][i];  
    }  
    B[i] = sum;  
}
```

S136\_1

```
for (int i=0;i<LEN;i++){  
    sum[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum[i] += A[j][i];  
    }  
    B[i]=sum[i];  
}
```

S136\_2

```
for (int i=0;i<LEN;i++)  
    B[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        B[i] += A[j][i];  
    }
```

S136

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Vectorization possible but seems inefficient  
**Exec. Time scalar code:** 3.7  
**Exec. Time vector code:** --  
**Speedup:** --



S136\_1

**Intel Nehalem**  
**report:** Permuted loop was vectorized.  
**scalar code:** 1.6  
**vector code:** 0.6  
**Speedup:** 2.6

S136\_2

**Intel Nehalem**  
**report:** Permuted loop was vectorized.  
**scalar code:** 1.6  
**vector code:** 0.6  
**Speedup:** 2.6

# Non-unit Stride – Example II

S136

```
for (int i=0;i<LEN;i++){  
    sum = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum += A[j][i];  
    }  
    B[i] = sum;  
}
```

S136\_1

```
for (int i=0;i<LEN;i++){  
    sum[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        sum[i] += A[j][i];  
    }  
    B[i]=sum[i];  
}
```

S136\_2

```
for (int i=0;i<LEN;i++)  
    B[i] = (float) 0.0;  
    for (int j=0;j<LEN;j++){  
        B[i] += A[j][i];  
    }
```

S136

**IBM Power 7**  
**Compiler report:** Loop was  
not SIMD vectorized  
**Exec. Time scalar code:** 2.0  
**Exec. Time vector code:** --  
**Speedup:** --

S136\_1

**IBM Power 7**  
**report:** Loop  
interchanging applied.  
Loop was SIMD  
vectorized  
**scalar code:** 0.4  
**vector code:** 0.2  
**Speedup:** 2.0

S136\_2

**IBM Power 7**  
**report:** Loop  
interchanging applied.  
Loop was SIMD  
**scalar code:** 0.4  
**vector code:** 0.16  
**Speedup:** 2.7



# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
  - Data Alignment
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics



# Conditional Statements – I

- Loops with conditions need #pragma vector always
  - Since the compiler does not know if vectorization will be profitable
  - The condition may prevent from an exception

```
#pragma vector always
for (int i = 0; i < LEN; i++){
    if (c[i] < (float) 0.0)
        a[i] = a[i] * b[i] + d[i];
}
```



# Conditional Statements – I

S137

```
for (int i = 0; i < LEN; i++){
    if (c[i] < (float) 0.0)
        a[i] = a[i] * b[i] + d[i];
}
```

S137\_1

```
#pragma vector always
for (int i = 0; i < LEN; i++){
    if (c[i] < (float) 0.0)
        a[i] = a[i] * b[i] + d[i];
}
```

S137

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Condition may protect exception  
**Exec. Time scalar code:** 10.4  
**Exec. Time vector code:** --  
**Speedup:** --

S137\_1

**Intel Nehalem**  
**Compiler report:** Loop was vectorized.  
**Exec. Time scalar code:** 10.4  
**Exec. Time vector code:** 5.0  
**Speedup:** 2.0



# Conditional Statements – I

S137

```
for (int i = 0; i < LEN; i++){
    if (c[i] < (float) 0.0)
        a[i] = a[i] * b[i] + d[i];
}
```

S137

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 4.0  
**Exec. Time vector code:** 1.5  
**Speedup:** 2.5

S137\_1

```
for (int i = 0; i < LEN; i++){
    if (c[i] < (float) 0.0)
        a[i] = a[i] * b[i] + d[i];
}
```

compiled with flag -qdebug=alwaysSPEC

S137\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 4.0  
**Exec. Time vector code:** 1.5  
**Speedup:** 2.5



# Conditional Statements

- Compiler removes *if conditions* when generating vector code

```
for (int i = 0; i < LEN; i++){  
    if (c[i] < (float) 0.0)  
        a[i] = a[i] * b[i] + d[i];  
}
```



# Conditional Statements

```
for (int i=0;i<1024;i++){
    if (c[i] < (float) 0.0)
        a[i]=a[i]*b[i]+d[i];
}
```

rC	2	-1	1	-2
rCmp	False	True	False	True
rThen	0	3.2	0	3.2
rElse	1.	0	1.	0
rS	1.	3.2	1.	3.2

```
vector bool char = rCmp
vector float r0={0.,0.,0.,0.};
vector float rA,rB,rC,rD,rS, rT,
rThen,rElse;
for (int i=0;i<1024;i+=4){
    // Load rA, rB, and rD;
    rCmp = vec_cmplt(rC, r0);
    rT= rA*rB+rD;
    rThen = vec_and(rT.rCmp);
    rElse = vec_andc(rA.rCmp);
    rs = vec_or(rthen, relse);
    //store rs
}
```



# Conditional Statements

```
for (int i=0;i<1024;i++){
    if (c[i] < (float) 0.0)
        a[i]=a[i]*b[i]+d[i];
}
```

Speedups will depend on the values on c[i]

Compiler tends to be conservative, as the condition may prevent from segmentation faults

```
vector bool char = rCmp
vector float r0={0.,0.,0.,0.};
vector float rA,rB,rC,rD,rS, rT,
rThen,rElse;
for (int i=0;i<1024;i+=4){
    // Load rA, rB, and rD;
    rCmp = vec_cmplt(rC, r0);
    rT= rA*rB+rD;
    rThen = vec_and(rT.rCmp);
    rElse = vec_andc(rA.rCmp);
    rs = vec_or(rthen, relse);
    //store rs
}
```



# Compiler Directives

- Compiler vectorizes many loops, but many more can be vectorized if the appropriate directives are used

Compiler Hints for Intel ICC	Semantics
#pragma ivdep	Ignore assume data dependences
#pragma vector always	override efficiency heuristics
#pragma novector	disable vectorization
__restrict__	assert exclusive access through pointer
__attribute__((aligned(int-val)))	request memory alignment
memalign(int-val,size);	malloc aligned memory
__assume_aligned(exp, int-val)	assert alignment property



# Compiler Directives

- Compiler vectorizes many loops, but many more can be vectorized if the appropriate directives are used

Compiler Hints for IBM XLC	Semantics
#pragma ibm independent_loop	Ignore assumed data dependences
#pragma nosimd	disable vectorization
__restrict__	assert exclusive access through pointer
__attribute__((aligned(int-val)))	request memory alignment
memalign(int-val,size);	malloc aligned memory
__alignx (int-val, exp)	assert alignment property



# Outline

1. Intro
2. Data Dependences (Definition)
3. Overcoming limitations to SIMD-Vectorization
  - Data Dependences
  - Data Alignment
  - Aliasing
  - Non-unit strides
  - Conditional Statements
4. Vectorization with intrinsics



# Access the SIMD through intrinsics

- Intrinsics are vendor/architecture specific
- We will focus on the Intel vector intrinsics
- Intrinsics are useful when
  - the compiler fails to vectorize
  - when the programmer thinks it is possible to generate better code than the one produced by the compiler



# The Intel SSE intrinsics Header file

- SSE can be accessed using intrinsics.
- You must use one of the following header files:
  - `#include <xmmmintrin.h>` (for SSE)
  - `#include <emmintrin.h>` (for SSE2)
  - `#include <pmmmintrin.h>` (for SSE3)
  - `#include <smmmintrin.h>` (for SSE4)
- These include the prototypes of the intrinsics.



# Intel SSE intrinsics Data types

- We will use the following data types:
  - `__m128` packed single precision (vector XMM register)
  - `__m128d` packed double precision (vector XMM register)
  - `__m128i` packed integer (vector XMM register)
- Example

```
#include <xmmmintrin.h>
int main ( ) {
    ...
    __m128 A, B, C; /* three packed s.p. variables */
    ...
}
```



# Intel SSE intrinsic Instructions

- Intrinsics operate on these types and have the format:  
`_mm_instruction_suffix(...)`
- Suffix can take many forms. Among them:
  - `ss` scalar single precision
  - `ps` packed (vector) single precision
  - `sd` scalar double precision
  - `pd` packed double precision
  - `si#` scalar integer (8, 16, 32, 64, 128 bits)
  - `su#` scalar unsigned integer (8, 16, 32, 64, 128 bits)



# Intel SSE intrinsics Instructions – Examples

- Load four 16-byte aligned single precision values in a vector:

```
float a[4]={1.0,2.0,3.0,4.0}//a must be 16-byte aligned
__m128 x = _mm_load_ps(a);
```

- Add two vectors containing four single precision values:

```
__m128 a, b;
__m128 c = _mm_add_ps(a, b);
```



# Intrinsics (SSE)

```
#define n 1024
__attribute__((aligned(16)))
float a[n], b[n], c[n];

int main() {
for (i = 0; i < n; i++) {
    c[i]=a[i]*b[i];
}
}
```

```
#include <xmmmintrin.h>
#define n 1024
__attribute__((aligned(16))) float
a[n], b[n], c[n];

int main() {
__m128 rA, rB, rC;
for (i = 0; i < n; i+=4) {
    rA = _mm_load_ps(&a[i]);
    rB = _mm_load_ps(&b[i]);
    rC= _mm_mul_ps(rA,rB);
    _mm_store_ps(&c[i], rC);
}}
```



# Intel SSE intrinsics

## A complete example

```
#define n 1024
```



```
int main() {
float a[n], b[n], c[n];
for (i = 0; i < n; i+=4) {
    c[i:i+3]=a[i:i+3]+b[i:i+3];
}
}
```

```
#include <xmmmintrin.h>
#define n 1024
__attribute__((aligned(16))) float
a[n], b[n], c[n];

int main() {
__m128 rA, rB, rC;
for (i = 0; i < n; i+=4) {
    rA = _mm_load_ps(&a[i]);
    rB = _mm_load_ps(&b[i]);
    rC= _mm_mul_ps(rA,rB);
    _mm_store_ps(&c[i], rC);
}}
```



# Intel SSE intrinsics

## A complete example

```
#define n 1024

int main() {
    float a[n], b[n], c[n];
    for (i = 0; i < n; i+=4) {
        c[i:i+3]=a[i:i+3]
    }
}
```

Declare 3 vector  
registers

```
#include <xmmmintrin.h>
#define n 1024
__attribute__((aligned(16))) float
a[n], b[n], c[n];

int main() {
    __m128 rA, rB, rC;
    for (i = 0; i < n; i+=4) {
        rA = _mm_load_ps(&a[i]);
        rB = _mm_load_ps(&b[i]);
        rC= _mm_mul_ps(rA,rB);
        _mm_store_ps(&c[i], rC);
    }
}
```



# Intel SSE intrinsics

## A complete example

```
#define n 1000

int main() {
    float a[n], b[n], c[n];
    for (i = 0; i < n; i+=4) {
        c[i:i+3]=a[i:i+3]+b[i:i+3];
    }
}
```

Execute vector statements

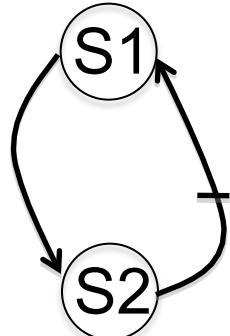
```
#include <xmmmintrin.h>
#define n 1024
__attribute__((aligned(16))) float
a[n], b[n], c[n];

int main() {
    __m128 rA, rB, rC;
    for (i = 0; i < n; i+=4) {
        rA = _mm_load_ps(&a[i]);
        rB = _mm_load_ps(&b[i]);
        rC= _mm_mul_ps(rA, rB);
        _mm_store_ps(&c[i], rC);
    }
}
```

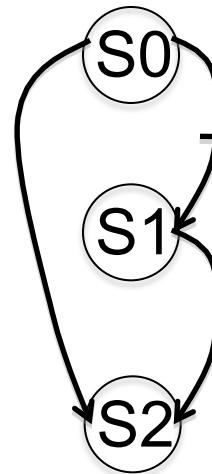


# Node Splitting

```
for (int i=0;i<LEN-1;i++){  
S1 a[i]=b[i]+c[i];  
S2 d[i]=(a[i]+a[i+1])*(float)0.5;  
}
```



```
for (int i=0;i<LEN-1;i++){  
S0 temp[i]=a[i+1];  
S1 a[i]=b[i]+c[i];  
S2 d[i]=(a[i]+temp[i])*(float) 0.5  
}
```



# Node Splitting with intrinsics

```
for (int i=0;i<LEN-1;i++){
    a[i]=b[i]+c[i];
    d[i]=(a[i]+a[i+1])*(float)0.5;
}
```

```
for (int i=0;i<LEN-1;i++){
    temp[i]=a[i+1];
    a[i]=b[i]+c[i];
    d[i]=(a[i]+temp[i])*(float)0.5;
}
```

Which code runs faster ?

Why?

```
#include <xmmmintrin.h>
#define n 1000

int main() {
    __m128 rA1, rA2, rB, rC, rD;
    __m128 r5=_mm_set1_ps((float)0.5)
    for (i = 0; i < LEN-4; i+=4) {
        rA2= _mm_loadu_ps(&a[i+1]);
        rB= _mm_load_ps(&b[i]);
        rC= _mm_load_ps(&c[i]);
        rA1= _mm_add_ps(rB, rC);
        rD= _mm_mul_ps(_mm_add_ps(rA1,rA2),r5);
        _mm_store_ps(&a[i], rA1);
        _mm_store_ps(&d[i], rD);
    }
}
```



# Node Splitting with intrinsics

S126

```
for (int i=0;i<LEN-1;i++){
    a[i]=b[i]+c[i];
    d[i]=(a[i]+a[i+1])*(float)0.5;
}
```

S126\_1

```
for (int i=0;i<LEN-1;i++){
    temp[i]=a[i+1];
    a[i]=b[i]+c[i];
    d[i]=(a[i]+temp[i])*(float)0.5;
}
```

S126\_2

```
#include <xmmmintrin.h>
#define n 1000

int main() {
    __m128 rA1, rA2, rB, rC, rD;
    __m128 r5=_mm_set1_ps((float)0.5)
    for (i = 0; i < LEN-4; i+=4) {
        rA2= _mm_loadu_ps(&a[i+1]);
        rB= _mm_load_ps(&b[i]);
        rC= _mm_load_ps(&c[i]);
        rA1= _mm_add_ps(rB, rC);
        rD= _mm_mul_ps(_mm_add_ps(rA1,rA2),r5);
        _mm_store_ps(&a[i], rA1);
        _mm_store_ps(&d[i], rD);
    }
}
```



# Node Splitting with intrinsics

S126

**Intel Nehalem**  
**Compiler report:** Loop was not vectorized. Existence of vector dependence  
**Exec. Time scalar code:** 12.6  
**Exec. Time vector code:** --  
**Speedup:** --

S126\_1

**Intel Nehalem**  
**Compiler report:** Loop was vectorized.  
**Exec. Time scalar code:** 13.2  
**Exec. Time vector code:** 9.7  
**Speedup:** 1.3

S126\_2

**Intel Nehalem**  
**Exec. Time intrinsics:** 6.1  
**Speedup (versus vector code):** 1.6



# Node Splitting with intrinsics

S126

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 3.8  
**Exec. Time vector code:** 1.7  
**Speedup:** 2.2

S126\_1

**IBM Power 7**  
**Compiler report:** Loop was SIMD  
vectorized  
**Exec. Time scalar code:** 5.1  
**Exec. Time vector code:** 2.4  
**Speedup:** 2.0

S126\_2

**IBM Power 7**  
**Exec. Time intrinsics:** 1.6  
**Speedup (versus vector code):** 1.5



# Summary

- Microprocessor vector extensions can contribute to improve program performance and the amount of this contribution is likely to increase in the future as vector lengths grow.
- Compilers are only partially successful at vectorizing
- When the compiler fails, programmers can
  - add compiler directives
  - apply loop transformations
- If after transforming the code, the compiler still fails to vectorize (or the performance of the generated code is poor), the only option is to program the vector extensions directly using intrinsics or assembly language.



# Data Dependencies

- The correctness of many many loop transformations including vectorization can be decided using dependences.
- A good introduction to the notion of dependence and its applications can be found in D. Kuck, R. Kuhn, D. Padua, B. Leasure, M. Wolfe: Dependence Graphs and Compiler Optimizations. POPL 1981.



# Compiler Optimizations

- For a longer discussion see:
  - Kennedy, K. and Allen, J. R. 2002 Optimizing Compilers for Modern Architectures: a Dependence-Based Approach. Morgan Kaufmann Publishers Inc.
  - U. Banerjee. Dependence Analysis for Supercomputing. Kluwer Academic Publishers, Norwell, Mass., 1988.
  - Advanced Compiler Optimizations for Supercomputers, by David Padua and Michael Wolfe in Communications of the ACM, December 1986, Volume 29, Number 12.
  - Compiler Transformations for High-Performance Computing, by David Bacon, Susan Graham and Oliver Sharp, in ACM Computing Surveys, Vol. 26, No. 4, December 1994.



# Algorithms

- W. Daniel Hillis and Guy L. Steele, Jr.. 1986.  
Data parallel algorithms. *Commun. ACM* 29, 12  
(December 1986), 1170-1183.
- Shyh-Ching Chen, D.J. Kuck, "Time and Parallel  
Processor Bounds for Linear Recurrence  
Systems," IEEE Transactions on Computers, pp.  
701-717, July, 1975



# Thank you

## Questions?

María Garzarán, Saeed Maleki  
William Gropp and David Padua  
[{garzaran,maleki,wgropp,padua}@illinois.edu](mailto:{garzaran,maleki,wgropp,padua}@illinois.edu)



# Program Optimization Through Loop Vectorization

María Garzarán, Saeed Maleki

William Gropp and David Padua

{garzaran,maleki,wgropp,padua}@illinois.edu

*Department of Computer Science*

*University of Illinois at Urbana-Champaign*



# Back-up Slides



# Measuring execution time

```
time1 = time();  
  
for (i=0; i<32000; i++)  
    c[i] = a[i] + b[i];
```

```
time2 = time();
```



# Measuring execution time

- Added an outer loop that runs (serially)
  - to increase the running time of the loop

```
time1 = time();
for (j=0; j<200000; j++){
    for (i=0; i<32000; i++)
        c[i] = a[i] + b[i];

}
time2 = time();
```



# Measuring execution times

- Added an outer loop that runs (serially)
  - to increase the running time of the loop
- **Call a dummy () function that is compiled separately**
  - to avoid loop interchange or dead code elimination

```
time1 = time();
for (j=0; j<200000; j++){
    for (i=0; i<32000; i++)
        c[i] = a[i] + b[i];
    dummy();
}
time2 = time();
```



# Measuring execution times

- Added an outer loop that runs (serially)
  - to increase the running time of the loop
- Call a dummy () function that is compiled separately
  - to avoid loop interchange or dead code elimination
- **Access the elements of one output array and print the result**
  - to avoid dead code elimination

```
time1 = time();
for (j=0; j<200000; j++){
    for (i=0; i<32000; i++)
        c[i] = a[i] + b[i];
    dummy();
}
time2 = time();
for (j=0; j<32000; j++)
    ret+= a[i];
printf ("Time %f, result %f", (time2 -time1), ret);
```



# Compiling

- Intelicc scalar code  
`icc -O3 -no-vec dummy.o tsc.o -o runnovec`
- Intelicc vector code  
`icc -O3 -vec-report[n] -xSSE4.2 dummy.o tsc.o -o runvec`

[n] can be 0,1,2,3,4,5

- **vec-report0**, no report is generated
- **vec-report1**, indicates the line number of the loops that were vectorized
- **vec-report2 .. 5**, gives a more detailed report that includes the loops that were not vectorized and the reason for that.



# Compiling

```
flags = -O3 -qaltivec -qhot -qarch=pwr7 -qtune=pwr7  
-qipa=malloc16 -qdebug=NSIMDCOST  
-qdebug=alwayspec -qreport
```

- IBM xlc scalar code  
`xlc -qnoenablevmx dummy.o tsc.o -o runovec`
- IBM vector code  
`xlc -qenablevmx dummy.o tsc.o -o runvec`



# Strip Mining

This transformation improves locality and is usually combined with vectorization



# Strip Mining

```
for (i=1; i<LEN; i++) {  
    a[i] = b[i];  
    c[i] = c[i-1] + a[i];  
}
```

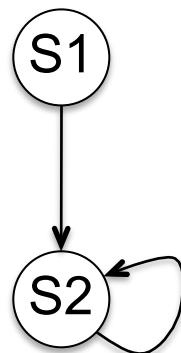
- first statement can be vectorized
- second statement cannot be vectorized because of self-true dependence



```
for (i=1; i<LEN; i++)  
    a[i] = b[i];  
  
for (i=1; i<LEN; i++)  
    c[i] = c[i-1] + a[i];
```

By applying loop distribution the compiler will vectorize the first statement

But, ... loop distribution will increase the cache miss ratio if array a[] is large



# Strip Mining

## Loop Distribution

```
for (i=1; i<LEN; i++)
    a[i] = b[i];
for (i=1; i<LEN; i++)
    c[i] = c[i-1] + a[i];
```

## Strip Mining

```
for (i=1; i<LEN; i
    +=strip_size){
    for (j=i; j<strip_size; j++)
        a[j] = b[j];
    for (j=i; j<strip_size; j++)
        c[j] = c[j-1] + a[j];
}
```

strip\_size is usually a small value (4, 8, 16 or 32).



# Strip Mining

- Another example

```
int v[N];
...
for (int i=0;i<N;i++){
    Transform (v[i]);
}
for (int i=0;i<N;i++){
    Light (v[i]);
}
```

```
int v[N];
...
for (int i=0;i<N;i+=strip_size){
    for (int j=i;j<strip_size;j++){
        Transform (v[j]);
    }
    for (int j=i;j<strip_size;j++){
        Light (v[j]);
    }
}
```

